


# Practical Deep Learning

Episode 0, 2025

## ML recap. Neural Nets 101



# ‘bout the course

## First module:

- Layers and optimization (Conv, Transformer, ...)
- Popular vision and NLP tasks

## Second module:

- Advanced stuff: LLM/Agents, Diffusion, RL, ...
- Elective topics (we'll run a poll!)

## Workload:

- One small assignment after each week's practice
- It's open-source! Go contribute :)

# Linear Regression

Model:

$$X \longrightarrow Wx + b \longrightarrow Y^{\text{pred}}$$

Objective function:

$$L = \sum_i (y_i - y_i^{\text{pred}})^2$$

Optimization (exact):

$$w = (X^T X)^{-1} X^T y$$

# Linear Regression

Model:

$$X \longrightarrow Wx + b \longrightarrow Y^{pred}$$

Objective function:

$$L = \sum_i (y_i - y_i^{pred})^2$$

Optimization (iterative):


$$w_0 \leftarrow 0$$

$$w_{i+1} \leftarrow w_i - \alpha \frac{\partial L}{\partial W}$$

$$\frac{\partial L}{\partial W} = \sum_i -2x(y_i - (wx_i + b))$$

# Logistic Regression

$$X \rightarrow Wx + b \rightarrow P(y)$$




$$P(y) = \sigma(Wx + b)$$

Objective function ?

# Logistic Regression

Model:

$$X \rightarrow Wx + b \rightarrow P(y|x)$$



Objective function:

$$L = - \sum_i y_i \log P(y|x_i) + (1 - y_i) \log (1 - P(y|x_i))$$

Optimization (iterative):

You guessed it!

# Logistic Regression

Model:

$$\begin{array}{ccc} a_{[y=a]} = W_a x + b_a & \xrightarrow{\quad} & P(y=a|X) \\ X \longrightarrow a_{[y=b]} = W_b x + b_b & \xrightarrow{\quad} & P(y=b|X) \\ a_{[y=c]} = W_c x + b_c & \xrightarrow{\quad} & P(y=c|X) \\ & \sum_j e^{a_{[y=j]}} & \end{array}$$

Objective function:


$$L = - \sum_i \log P(y_i^{correct} | x_i)$$

# Gradient descent


Update:

$$w_{i+1} \leftarrow w_i - \alpha \frac{\partial L}{\partial w}$$


- $\alpha$  – learning rate  $\alpha < < 1$
- $L$  – loss function



# Nonlinear dependencies



What we have



What we want


How to get that?

# Feature extraction

Loss, for example:

$$L = - \sum_i y_i \log P(y|x_i) + (1 - y_i) \log (1 - P(y|x_i))$$

Model:



Training:

$$\underset{\theta_1}{\operatorname{argmin}} L(y, P(y|x))$$




Features would tune to your problem automatically!

# What do we want, exactly?

Loss, for example:

$$L = - \sum_i y_i \log P(y|x_i) + (1 - y_i) \log (1 - P(y|x_i))$$

Model:



Training:

?


$$\underset{\theta_1}{\operatorname{argmin}} L(y, P(y|x))$$

# What do we want, exactly?

Loss, for example:

$$L = - \sum_i y_i \log P(y|x_i) + (1 - y_i) \log (1 - P(y|x_i))$$


Model:



Gradients:  $\underset{\theta_2}{\operatorname{argmin}} L(y, P(y|x))$      $\underset{\theta_1}{\operatorname{argmin}} L(y, P(y|x))$

# Try linear


Model:



$$h_j = \sum_i w_{ij}^h x_i + b_j^h \quad y_{pred} = \sigma \left( \sum_j w_j^o h_j + b^o \right)$$

# Try linear

Model:



$$h_j = \sum_i w_{ij}^h x_i + b_j^h \quad y_{pred} = \sigma \left( \sum_j w_j^o h_j + b^o \right)$$

Output:

$$P(y|x) = \sigma \left( \sum_j w_j^o \left( \sum_i w_{ij}^h x_i + b_j^h \right) + b^o \right)$$

Is it any better than logistic regression?

# Try linear


$$P(y|x) = \sigma\left(\sum_j w_j^o \left(\sum_i w_{ij}^h x_i + b_j^h\right) + b^o\right)$$

$$w'_i = \sum_j w_j^o w_{ij}^h \quad b' = \sum_j w_j^o b_j^h + b^o$$

$$P(y|x) = \sigma\left(\sum_i w'_i x_i + b'\right)$$

# Try linear

Model:



$$h_j = \sum_i w_{ij}^h x_i + b_j^h \quad y_{pred} = \sigma \left( \sum_j w_j^o h_j + b^o \right)$$


Output:

$$P(y|x) = \sigma \left( \sum_j w_j^o \left( \sum_i w_{ij}^h x_i + b_j^h \right) + b^o \right)$$

Is it any better than logistic regression?

# Nonlinearity

Model:




$$h_j = \sigma\left(\sum_{j \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h\right)$$

$$y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$

# Nonlinearity

Model:




$$h_j = \sigma\left(\sum_{i \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h\right) \quad y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$

Output:  $P(y|x) = \sigma\left(\sum_j w_j^o \sigma\left(\sum_i w_{ij}^h x_i + b_j^h\right) + b^o\right)$


# Nonlinearity

Model:




$$h_j = \sigma\left(\sum_{i \in \{1, 2, \dots, n\}} w_{ij}^h x_i + b_j^h\right)$$

$$y_{pred} = \sigma\left(\sum_j w_j^o h_j + b^o\right)$$




# Nonlinearity

- $f(a) = 1/(1+e^a)$
- $f(a) = \tanh(a)$
- $f(a) = \max(0,a)$
- $f(a) = \log(1+e^a)$



# Training neural nets

Model:



Output:  $P(y|x) = \sigma\left(\sum_j w_j^o \sigma\left(\sum_i w_{ij}^h x_i + b_j^h\right) + b^o\right)$

$$\partial E - \log P_w(y_i|x_i)$$

Training:  $w := w - \alpha \frac{x_i, y_i}{\partial w}$

# Part II: Backpropagation

# Backpropagation

**TL;DR:** backprop = chain rule\*


$$\frac{\partial f(g(x))}{\partial x} = \frac{\partial f(g(x))}{\partial g(x)} \cdot \frac{\partial g(x)}{\partial x}$$

# Backpropagation


**TL;DR:** backprop = chain rule\*

$$\frac{\partial f(g(x))}{\partial x} = \frac{\partial f(g(x))}{\partial g(x)} \cdot \frac{\partial g(x)}{\partial x}$$

\* g and x can be vectors/vectors/tensors




# Backpropagation




$$\frac{\partial L(\sigma(\text{linear}_{w2,b2}(\sigma(\text{linear}_{w1,b1}(x)))))}{\partial w1} = \dots$$

# Backpropagation




$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial \sigma} \cdot$$

# Backpropagation




$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial \sigma} \cdot \frac{\partial \sigma}{\partial \text{linear}_{w_2, b_2}}.$$

# Backpropagation



$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial \sigma} \cdot \frac{\partial \sigma}{\partial linear_{w_2, b_2}} \cdot \frac{\partial linear_{w_2, b_2}}{\partial \sigma}$$

# Backpropagation



$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial \sigma} \cdot \frac{\partial \sigma}{\partial \text{linear}_{w_2, b_2}} \cdot \frac{\partial \text{linear}_{w_2, b_2}}{\partial \sigma} \cdot \frac{\partial \sigma}{\partial \text{linear}_{w_1, b_1}} \cdot \frac{\partial \text{linear}_{w_1, b_1}}{\partial w_1}$$

# Matrix derivatives

Let's compute:

$$\frac{\partial L(X \times W + b)}{\partial X} = \frac{\partial L(X \times W + b)}{\partial [X \times W + b]} \times \boxed{\text{What?}}$$

Variable shapes:

$X$

[batch size, features]

$W$

[features, outputs]

$b$

[outputs]

$$\frac{\partial L(X \times W + b)}{\partial X}$$

[batch size, features]

$$\frac{\partial L(X \times W + b)}{X \times W + b}$$

[batch size, outputs]

# Matrix derivatives

Let's compute:

*Hint: 1. figure out scalar case,  
2. match shapes for matrices*

$$\frac{\partial L(X \times W + b)}{\partial X} = \frac{\partial L(X \times W + b)}{\partial [X \times W + b]} \times$$

What?

Variable shapes:

$X$

[batch size, features]

$W$

[features, outputs]

$b$

[outputs]

$$\frac{\partial L(X \times W + b)}{\partial X}$$

[batch size, features]

$$\frac{\partial L(X \times W + b)}{X \times W + b}$$

[batch size, outputs]

# Matrix derivatives

**Let's compute:**

$$\frac{\partial L(X \times W + b)}{\partial X} = \frac{\partial L(X \times W + b)}{\partial [X \times W + b]} \times W^T$$

**Variable shapes:**

$X$

[batch size, features]

$W$

[features, outputs]

$b$

[outputs]

$$\frac{\partial L(X \times W + b)}{\partial X}$$

[batch size, features]

$$\frac{\partial L(X \times W + b)}{X \times W + b}$$

[batch size, outputs]

# Matrix derivatives

Let's compute:

$$\frac{\partial L(X \times W + b)}{\partial W} =$$

what?

Variable shapes:

$X$

[batch size, features]

$W$

[features, outputs]

$b$

[outputs]

$$\frac{\partial L(X \times W + b)}{\partial X}$$

[batch size, features]

$$\frac{\partial L(X \times W + b)}{X \times W + b}$$

[batch size, outputs]

# Matrix derivatives

**Let's compute:**

$$\frac{\partial L(X \times W + b)}{\partial W} = X^T \times \frac{\partial L}{\partial [X \times W + b]}$$

**Variable shapes:**

$X$

[batch size, features]

$W$

[features, outputs]

$b$

[outputs]

$$\frac{\partial L(X \times W + b)}{\partial X}$$

[batch size, features]

$$\frac{\partial L(X \times W + b)}{X \times W + b}$$

[batch size, outputs]

# Matrix derivatives (formulae)

$$\frac{\partial \sum_i \log p(y_i|x_i, w)}{\partial w} = \frac{\sum_i \partial \log p(y_i|x_i, w)}{\partial w}$$

$$\frac{\partial L(X \times W + b)}{\partial X} = \frac{\partial L}{\partial [X \times W + b]} \times W^T$$


$$\frac{\partial L(X \times W + b)}{\partial W} = X^T \times \frac{\partial L}{\partial [X \times W + b]}$$

$$\frac{\partial L(\sigma(x))}{\partial x} = \frac{\partial L}{\partial \sigma(x)} \cdot [\sigma(x) \cdot (1 - \sigma(x))]$$

Works for any kind of x  
(scalar, vector, matrix, tensor)

# Back to neural networks


Model:




Training:




# Part III: Biological inspiration



# Biological inspiration





# Biological inspiration



# Not actual neurons :)


- Neurons react in “spikes”, not real numbers
- Neurons maintain/change their states over time
- No one knows for sure how they “train”
- Neuroglial cells are important  
But noone knows, why



# Connectionist phrasebook


- Layer – a building block for NNs :
  - “Dense layer”:  $f(x) = Wx+b$
  - “Nonlinearity layer”:  $f(x) = \sigma(x)$
  - Input layer, output layer
  - A few more we gonna cover later
- Activation – layer output
  - i.e. some intermediate signal in the NN
- Backpropagation – a fancy word for “chain rule”

# Connectionist phrasebook



- “Train it via backprop!”

# Connectionist phrasebook



How do we train it?


# Part IV: Optimization

# Gradient descent

Update:

$$w_{i+1} \leftarrow w_i - \alpha \frac{\partial L}{\partial w}$$

- $\alpha$  – learning rate  $\alpha < < 1$
- $L$  – loss function



Can we do better?


# Newton-Raphson

Parameter update

$$w_{i+1} \leftarrow w_i - \alpha H_L^{-1} \frac{\partial L}{\partial w}$$

Hessian:

$$\mathbf{H} = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \dots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \dots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}.$$



Red: Newton-Raphson  
Green: gradient descent

Any drawbacks?

# Stochastic gradient descent


Loss function is mean over all data samples.

Approximate with 1 or few random samples.

Update:

$$w_{i+1} \leftarrow w_i - \alpha E \frac{\partial L}{\partial w}$$

- E – expectation
- Learning rate should decrease



# SGD with momentum

Idea: move towards “overall gradient direction”,  
Not just current gradient.

$$w_0 \leftarrow 0; v_0 \leftarrow 0$$

$$v_{i+1} \leftarrow \alpha \frac{\partial L}{\partial w} + \mu v_i$$

$$w_{i+1} \leftarrow w_i - v_{i+1}$$

Helps for noisy gradient / canyon problem



# SGD with momentum

Idea: move towards “overall gradient direction”,  
Not just current gradient.

$$w_0 \leftarrow 0; v_0 \leftarrow 0$$

$$v_{i+1} \leftarrow \alpha \frac{\partial L}{\partial w} + \mu v_i$$

$$w_{i+1} \leftarrow w_i - v_{i+1}$$



# AdaGrad

Idea: decrease learning rate individually for each parameter in proportion to sum of it's gradients so far.

$$G_t = \sum_{\tau=1}^t \left[ \frac{\partial L}{\partial w} \right]^2$$

“Total update path length”  
(for each parameter)

$$w_{t+1} = w_t - \frac{\eta}{\sqrt{G_t + \epsilon}} \frac{\partial L}{\partial w}$$


# RMSProp

Idea: make sure all gradient steps have approximately same magnitude (by keeping moving average of magnitude)

$$ms_{t+1} = \gamma \cdot ms_t + (1 - \gamma) \left\| \frac{\partial L}{\partial w} \right\|^2$$

$$w_{t+1} = w_t - \frac{\eta}{\sqrt{ms + \epsilon}} \frac{\partial L}{\partial w}$$

# All together



# Adam Optimizer

<https://arxiv.org/abs/1412.6980>

Update rule:

$$w_t = w_{t-1} - \eta \frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}$$

# Adam Optimizer

<https://arxiv.org/abs/1412.6980>

Update rule:

$$w_t = w_{t-1} - \eta \frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}$$

$\approx$  momentum

$\approx$  rmsprop

# Adam Optimizer

<https://arxiv.org/abs/1412.6980>

Update rule:

$$w_t = w_{t-1} - \eta \frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}$$

≈ momentum

≈ rmsprop

Statistics:

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

(Optional) bias correction

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}$$

$$\hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

Moving averages of gradient and squared gradient.

# Newer stuff

Minor Adam improvements: **Adan**, **AdEMAMix**

Memory savings: Adafactor, Adam8bit, Adammini

Matrix-level loss geometry: **Muon**, **Shampoo**, **Soap**


Other: **LAMB**(large batch), **LION**(sign), **Prodigy**(tuning-free)

Recent optimizer comparisons (September 2025):

**Fantastic Pretraining Optimizers and Where to Find Them**


**Benchmarking Optimizers for LLM Pretraining**

# Initialization, symmetry problem




- Initialize with zeros  
 $W \leftarrow 0$
- What will the first step look like?

# Initialization, symmetry problem



- Break the symmetry!
- Initialize with random numbers!  
 $W \leftarrow N(0,0.01)$ ?  
 $W \leftarrow U(0,0.1)$ ?
- Can get a bit better for deep NNs

# Initialization, symmetry problem



- Break the symmetry!
- Initialize with random numbers!  
 $W \leftarrow N(0,0.01)$ ?  
 $W \leftarrow U(0,0.1)$ ?
- Can get a bit better for deep NNs

# Potential caveats?

- Hardcore overfitting
- No “golden standard” for architecture
- Computationally heavy

# Nuff

**Let's go implement that!**



# Short break

[playground.tensorflow.org](https://playground.tensorflow.org)