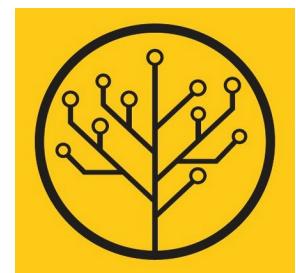


# Deep learning

Episode 2, 2025

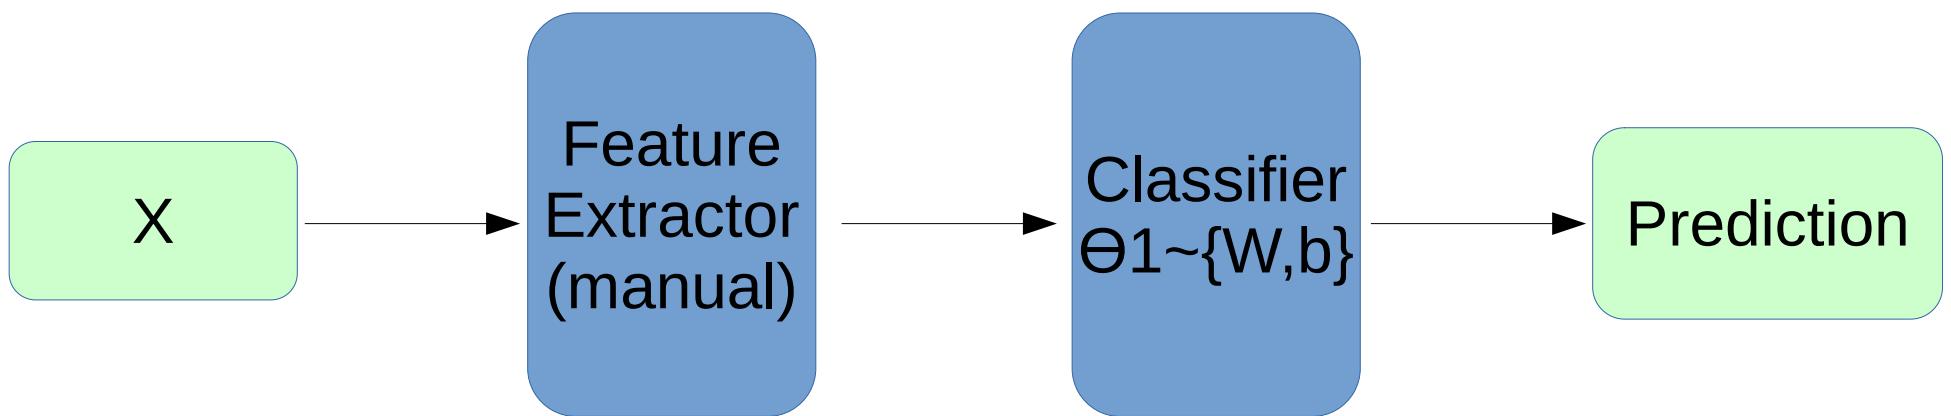
## Deep learning whereabouts

A catch-all lecture in philosophy,  
tricks and frameworks



Previously on deep learning...

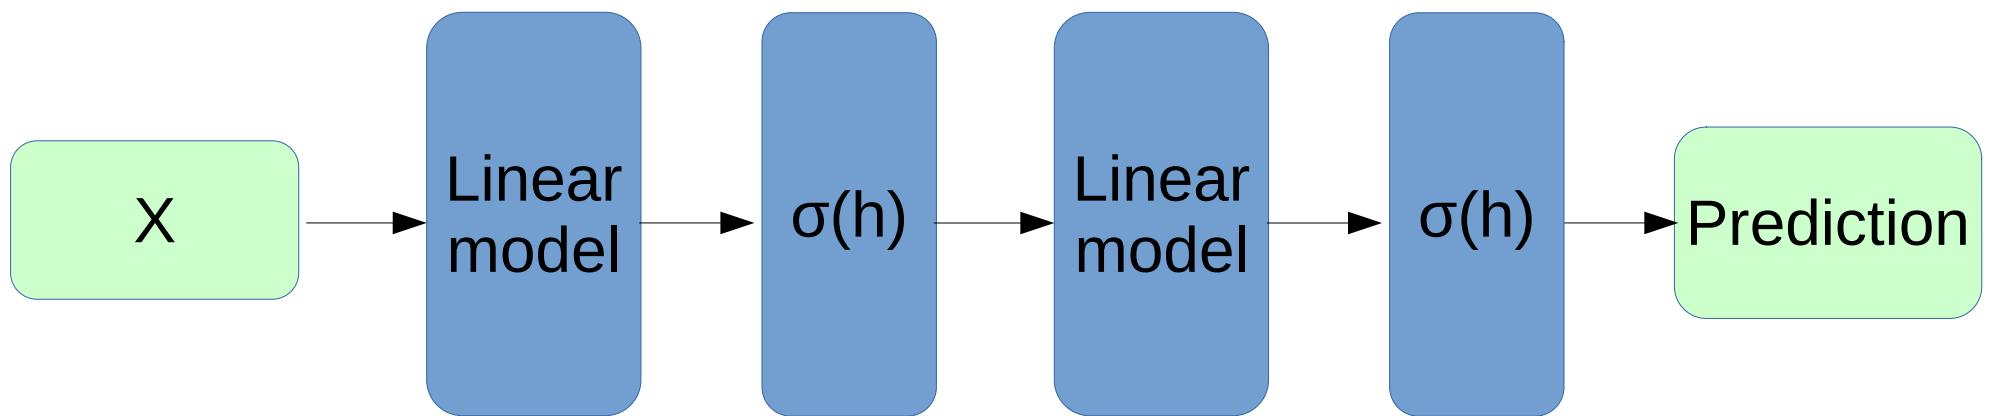
# Feature extraction





Features would tune to your problem automatically!

# Simple neural network



Trains with stochastic gradient descent!  
or momentum/rmsprop/adam/...

# Connectionist phrasebook

- Layer – a building block for NNs :
  - “Dense layer”:  $f(x) = Wx+b$
  - “Nonlinearity layer”:  $f(x) = \sigma(x)$
  - Input layer, output layer
  - A few more we gonna cover later
- Activation – layer output
  - i.e. some intermediate signal in the NN
- Backpropagation – a fancy word for “chain rule”

# Backpropagation

**TL;DR:** backprop = chain rule\*

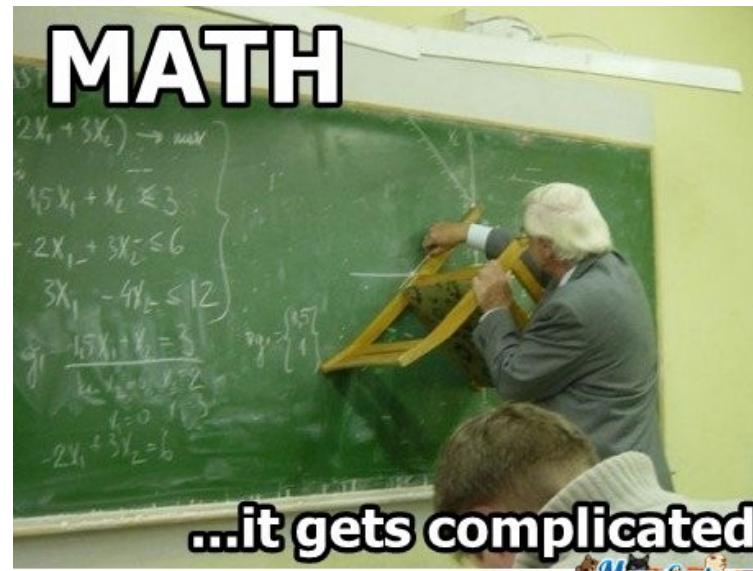
$$\frac{\partial f(g(x))}{\partial x} = \frac{\partial f(g(x))}{\partial g(x)} \cdot \frac{\partial g(x)}{\partial x}$$

# Backpropagation

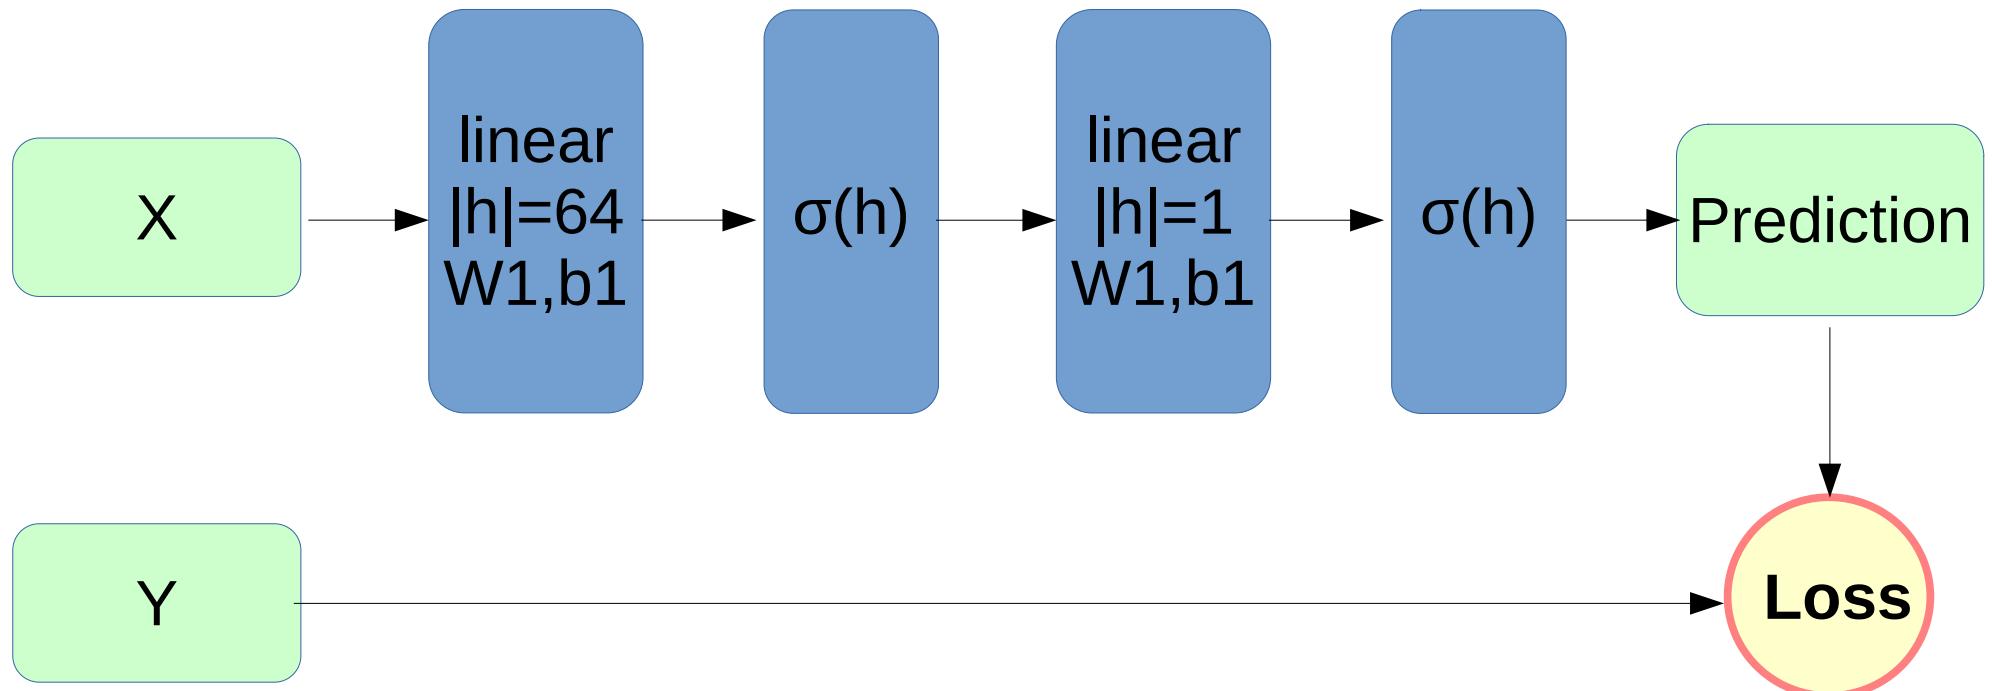
**TL;DR:** backprop = chain rule\*

$$\frac{\partial f(g(x))}{\partial x} = \frac{\partial f(g(x))}{\partial g(x)} \cdot \frac{\partial g(x)}{\partial x}$$

\* g and x can be vectors/vectors/tensors

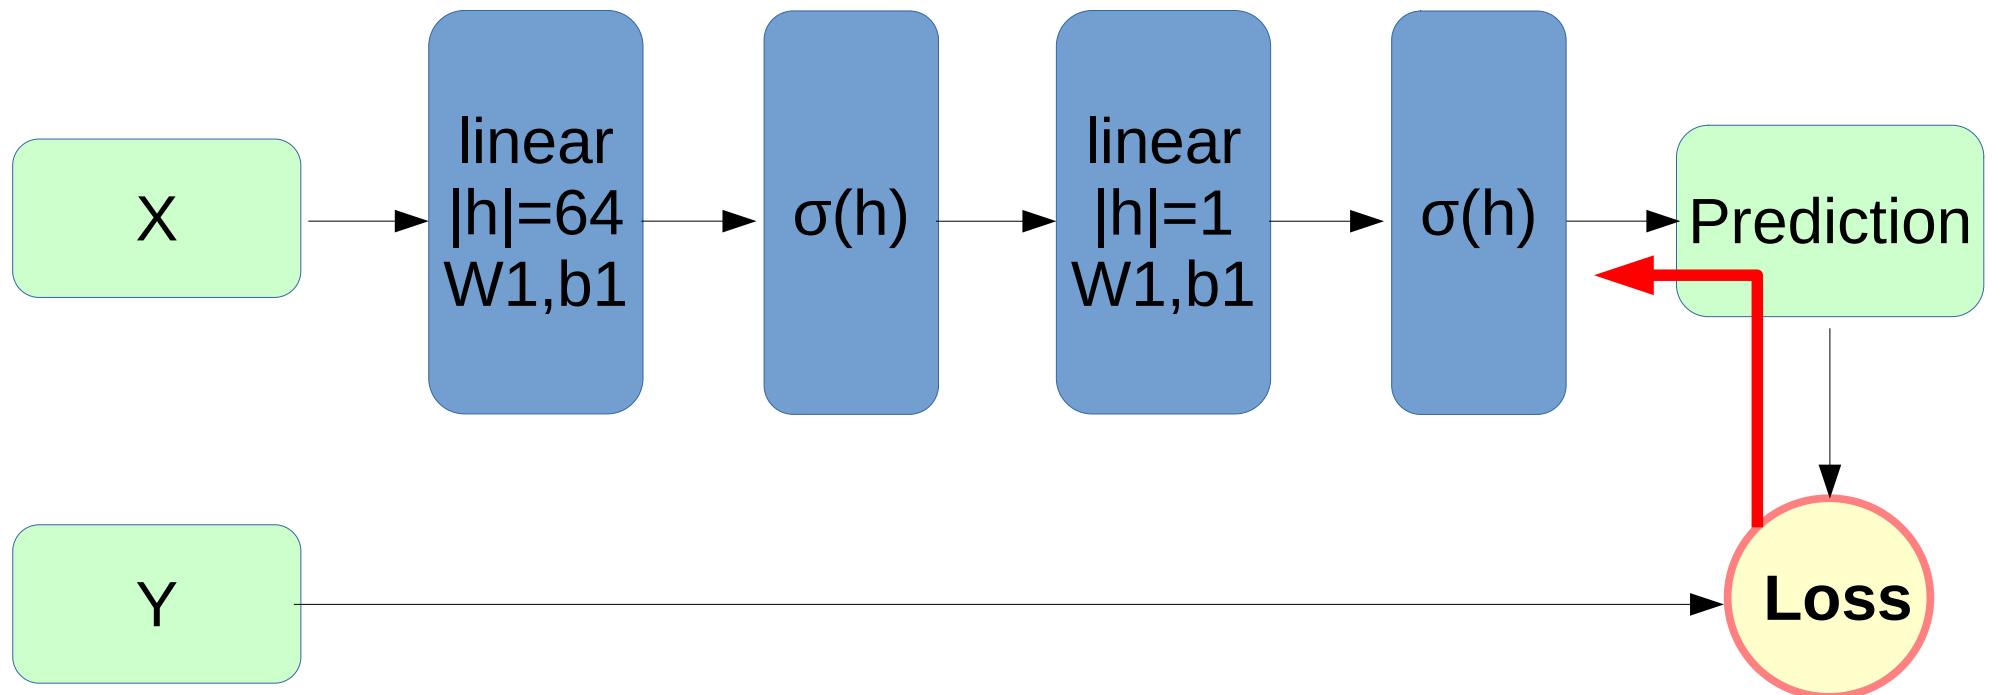


# Backpropagation



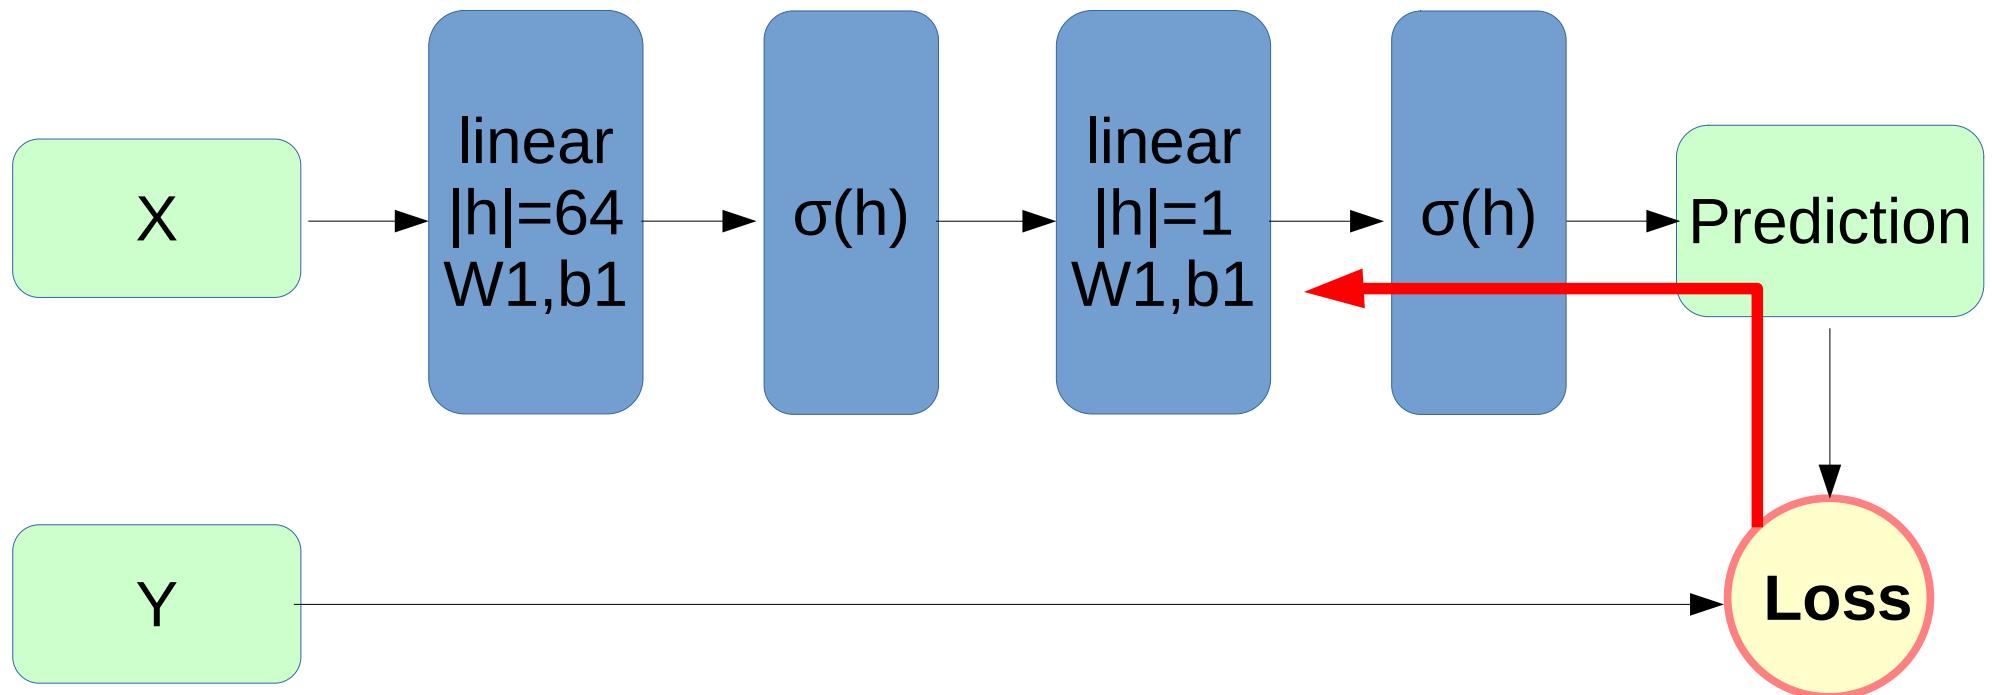
$$\frac{\partial L(\sigma(\text{linear}_{w2,b2}(\sigma(\text{linear}_{w1,b1}(x)))))}{\partial w1} = \dots$$

# Backpropagation



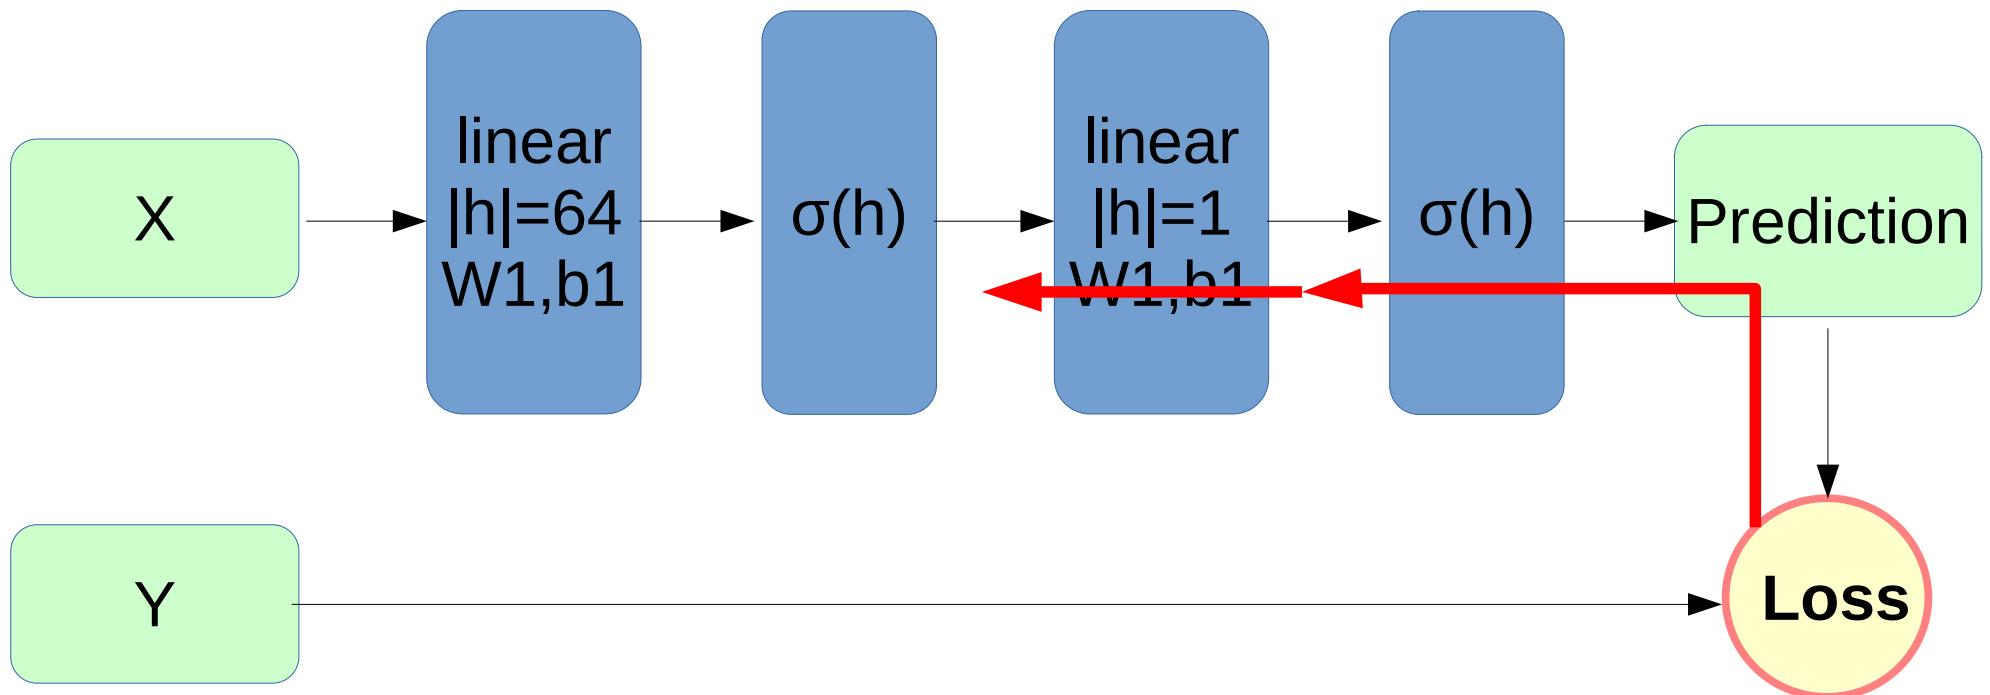
$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial \sigma} \cdot$$

# Backpropagation



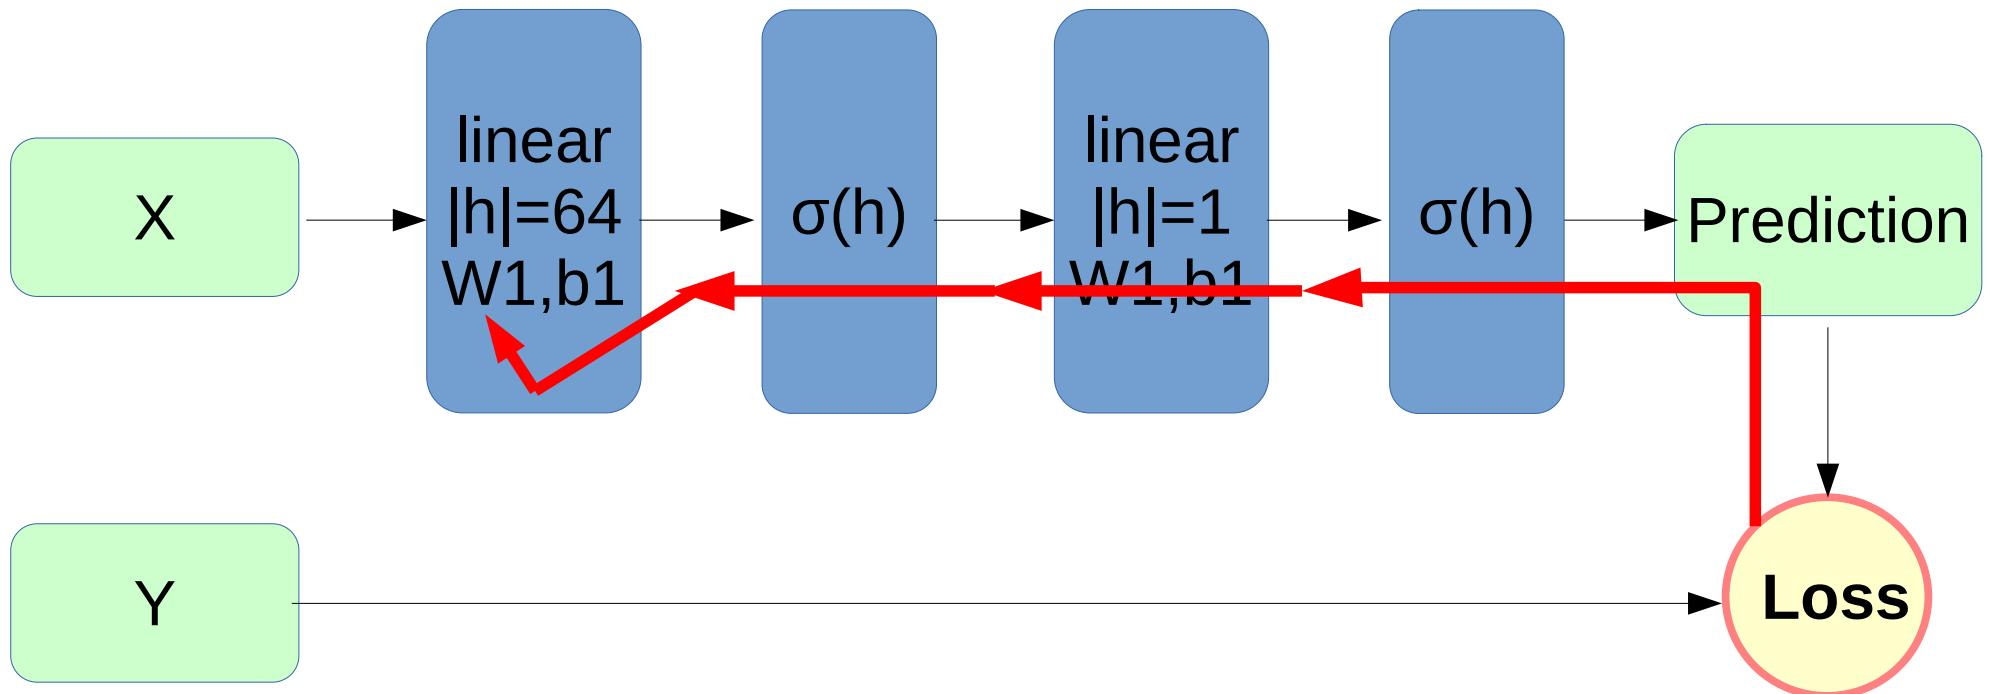
$$\frac{\partial L}{\partial w1} = \frac{\partial L}{\partial \sigma} \cdot \frac{\partial \sigma}{\partial linear_{w2, b2}}.$$

# Backpropagation



$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial \sigma} \cdot \frac{\partial \sigma}{\partial \text{linear}_{w_2, b_2}} \cdot \frac{\partial \text{linear}_{w_2, b_2}}{\partial \sigma}$$

# Backpropagation



$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial \sigma} \cdot \frac{\partial \sigma}{\partial linear_{w_2, b_2}} \cdot \frac{\partial linear_{w_2, b_2}}{\partial \sigma} \cdot \frac{\partial \sigma}{\partial linear_{w_1, b_1}} \cdot \frac{\partial linear_{w_1, b_1}}{\partial w_1}$$

# Matrix derivatives we used

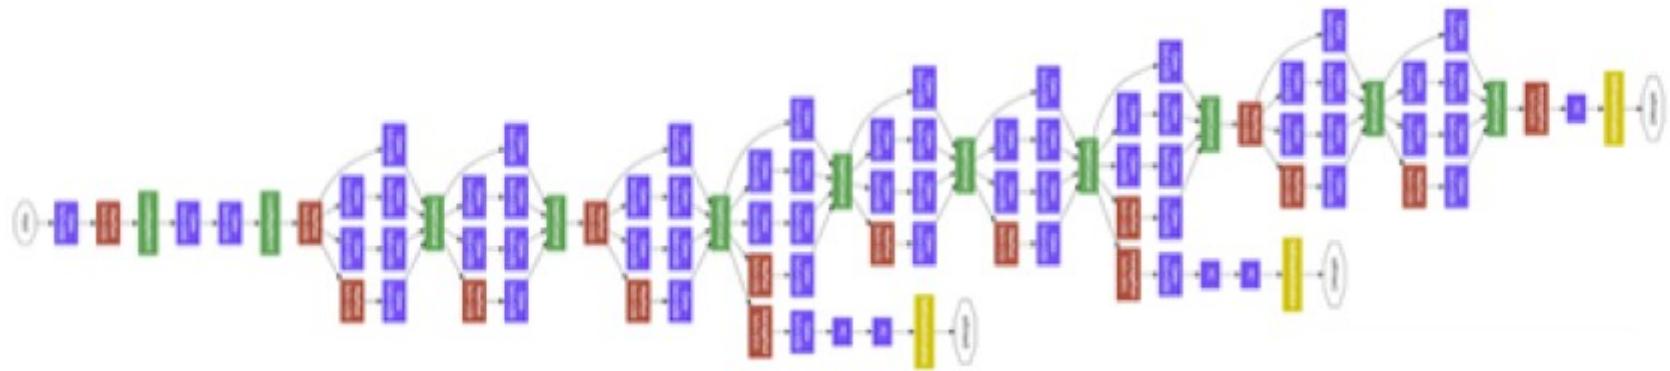
$$\text{sigmoid} : \frac{\partial L}{\partial \sigma(x)} \cdot [\sigma(x) \cdot (1 - \sigma(x))]$$

Works for any kind of  $x$   
(scalar, vector, matrix, tensor)

$$\text{linear over } X : \frac{\partial L}{\partial W \times X + b} \times W^T$$

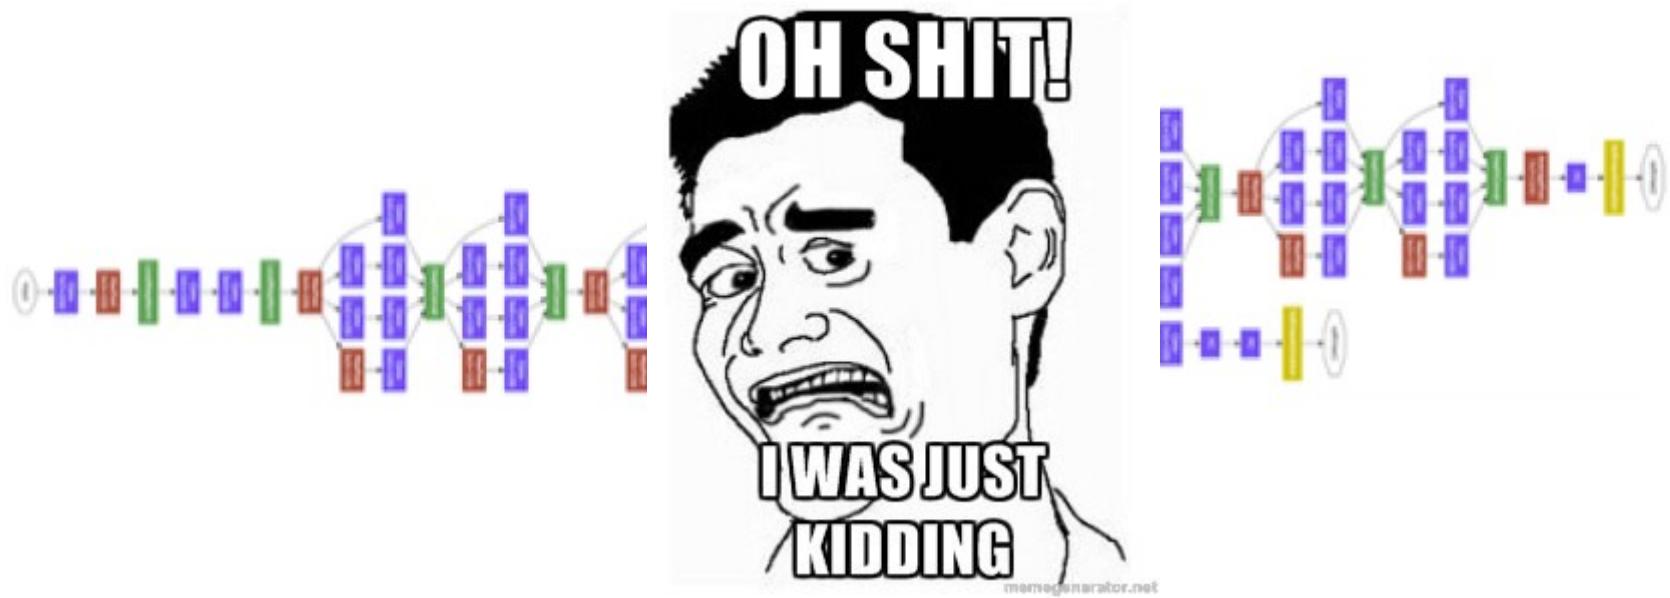
$$\text{linear over } W : \frac{1}{\|X\|} \cdot X^T \times \frac{\partial L}{\partial [X \times W + b]}$$

# And now let's differentiate



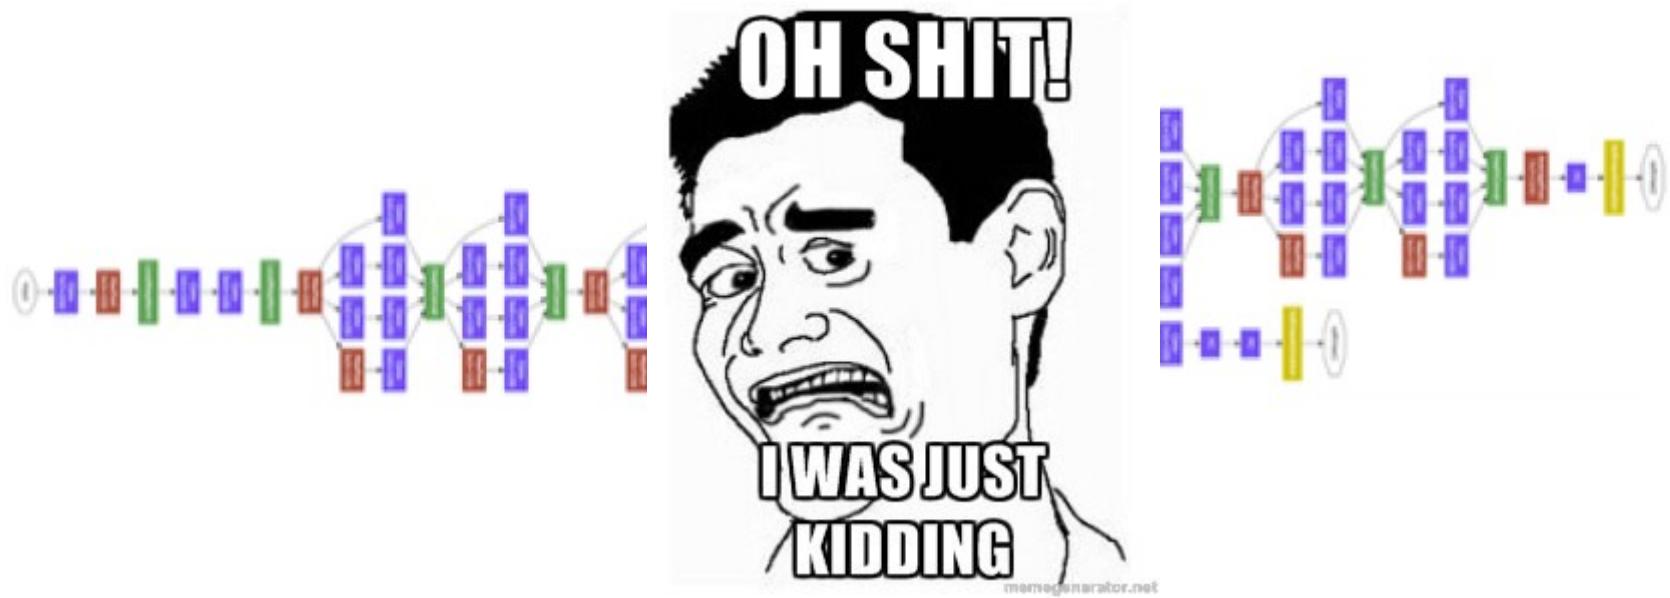
- 5+ types of layers
- each with different dimensions
- parallel branches with independent losses
- several nonlinearities

# And now let's differentiate



- 5+ types of layers
- each with different dimensions
- parallel branches with independent losses
- several nonlinearities

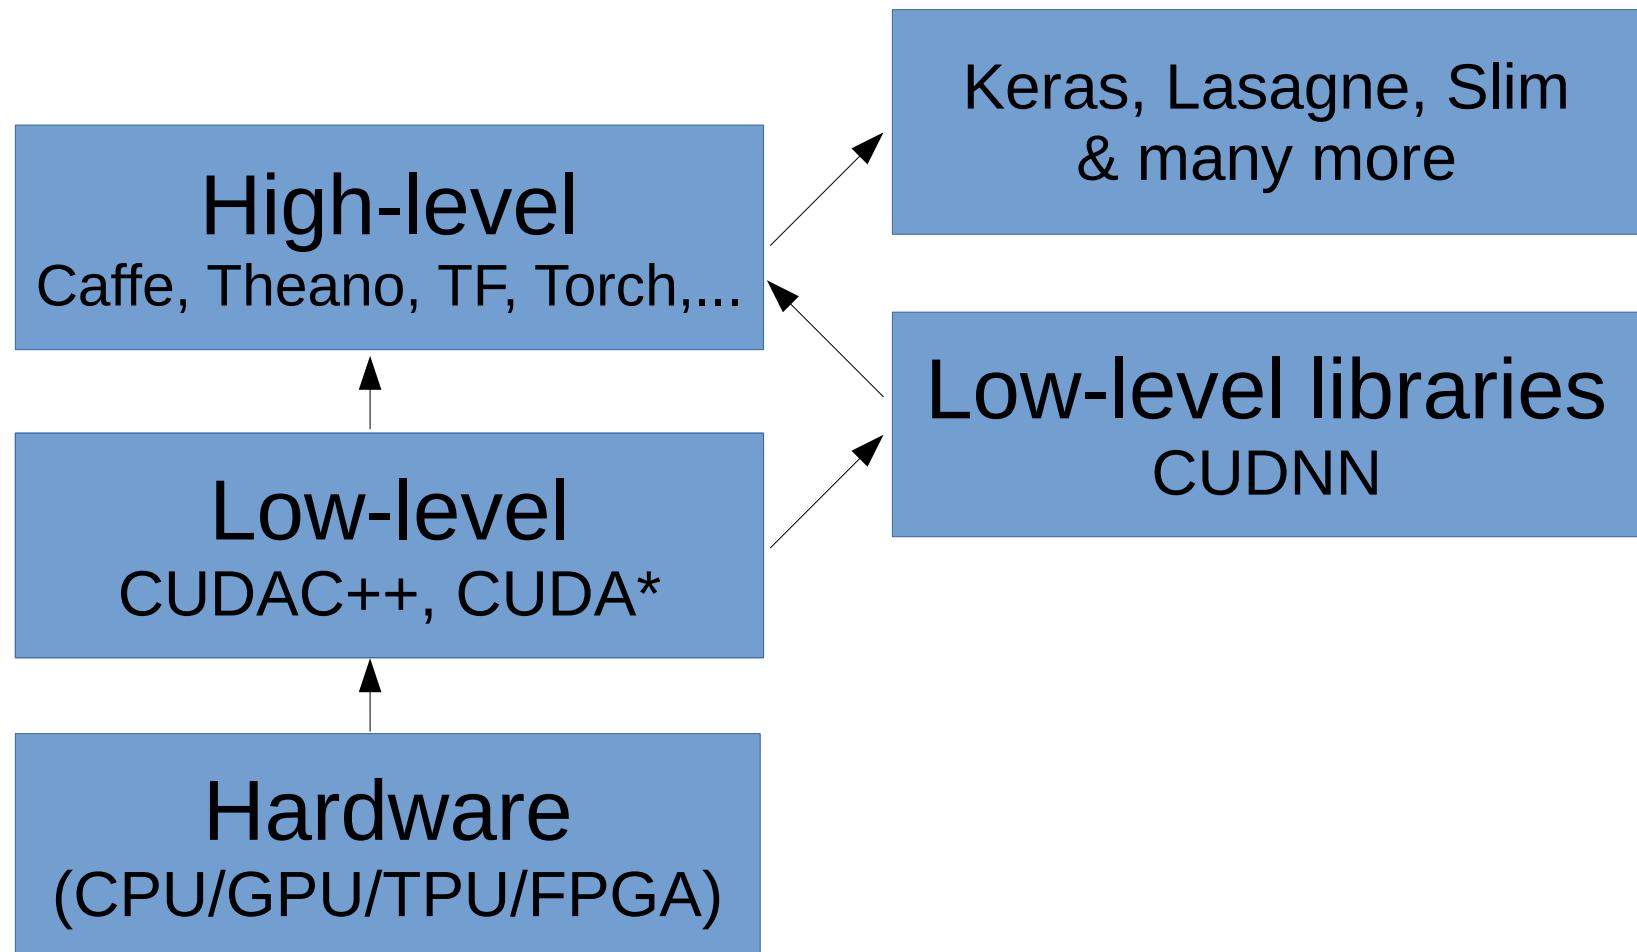
# Deep learning frameworks



- 5+ types of layers
- each with different dimensions
- parallel branches with independent losses
- several nonlinearities

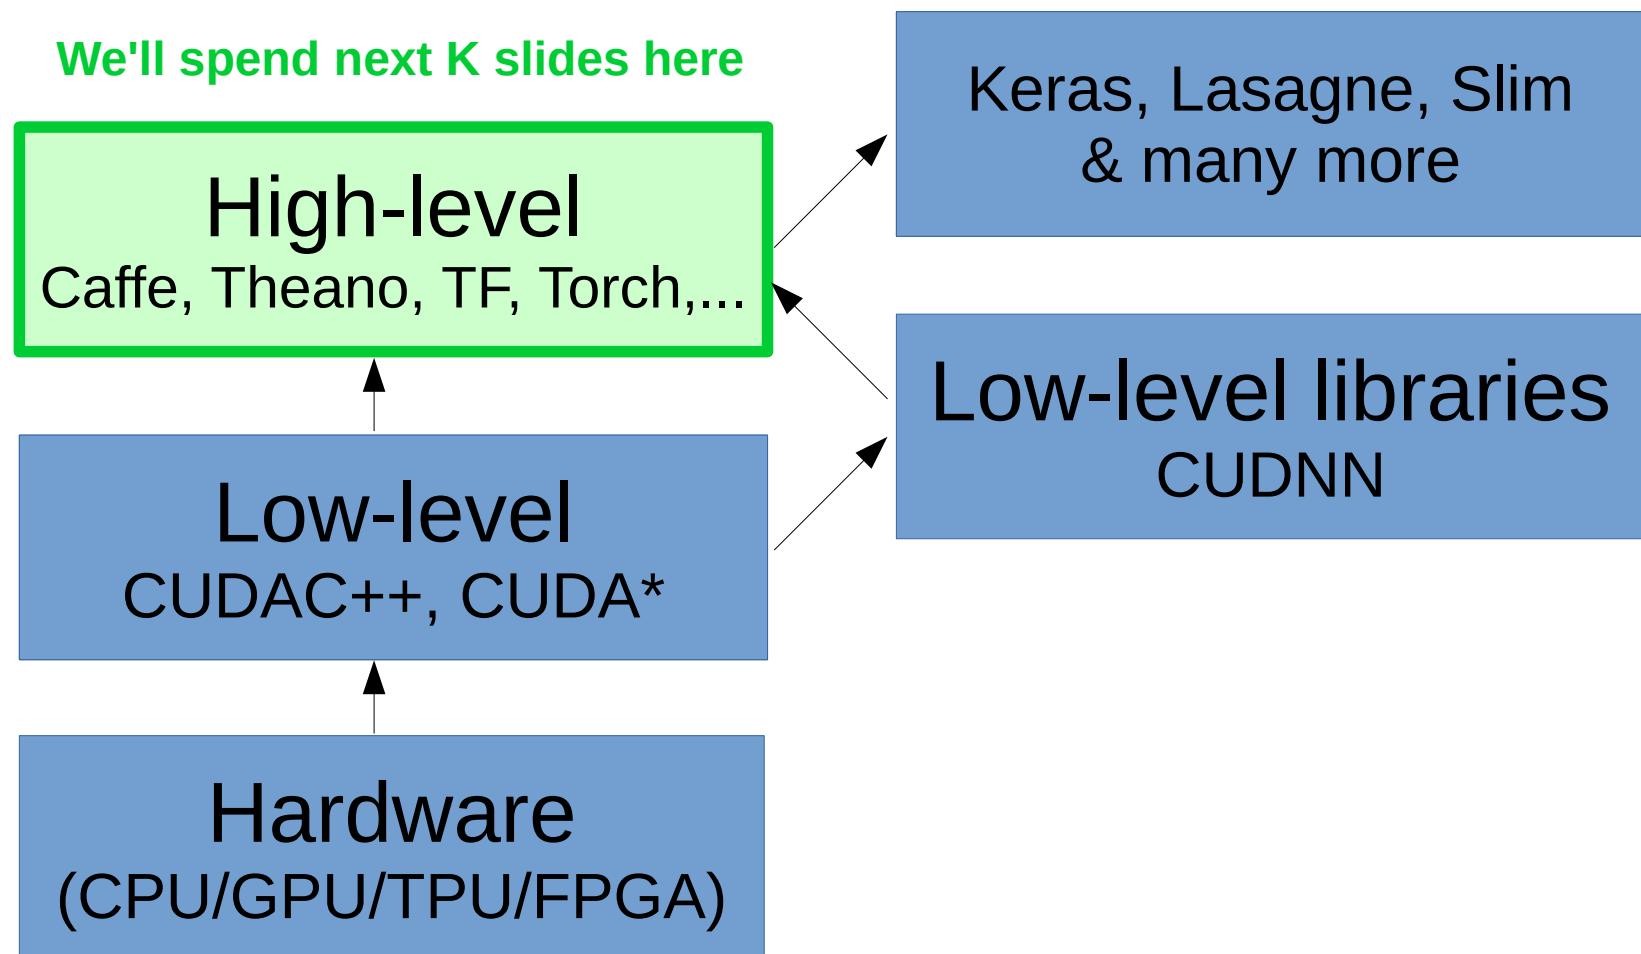
# Deep learning frameworks

- Core idea: helps you define and train neural nets



# Deep learning frameworks

- Core idea: helps you define and train neural nets



# Deep learning frameworks

Layer-based frameworks:

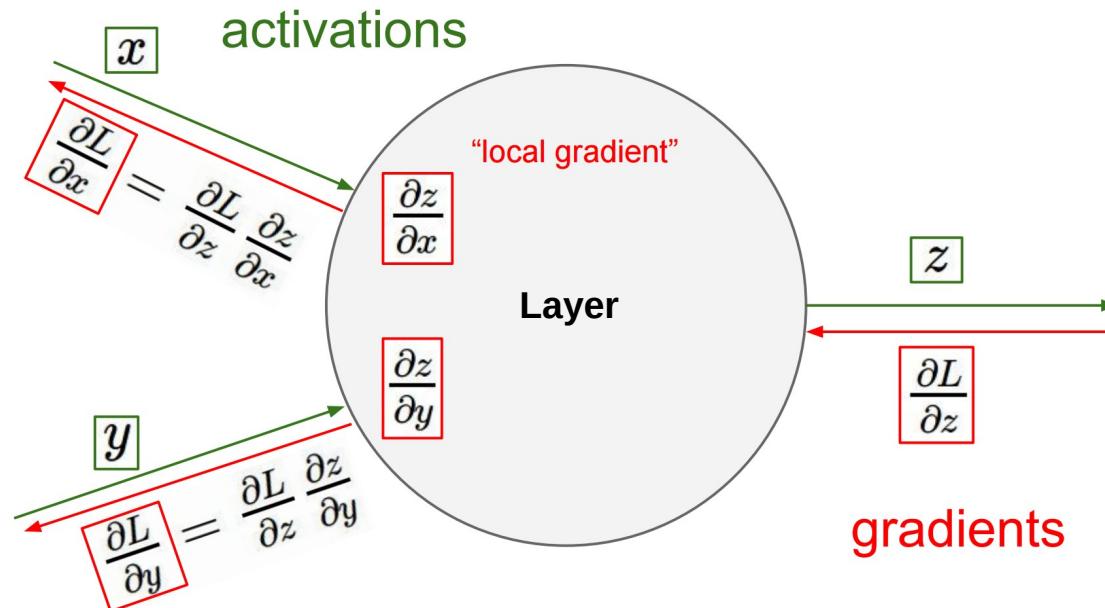
Same idea as in our hand-made neural net

# Deep learning frameworks

Layer-based frameworks:

Same idea as in our hand-made neural net

this one - <http://bit.ly/2w9kAHm>



# Deep learning frameworks

## Caffe

```
name: "LeNet"
layer {
    name: "conv1"
    type: "Convolution"
    bottom: "data"
    top: "conv1"
    param {lr_mult: 1}
    param {lr_mult: 2}
    convolution_param {
        num_output: 20
        kernel_size: 5
        stride: 1
        weight_filler {
            type: "xavier"
        }
    }
}
```

....

130 lines

You define model in config file  
by stacking layers.

Then train like this:

```
caffe train -solver
examples/mnist/lenet_solver.prototxt
```

# Deep learning frameworks

## Caffe

```
name: "LeNet"
layer {
    name: "conv1"
    type: "Convolution"
    bottom: "data"
    top: "conv1"
    param {lr_mult: 1}
    param {lr_mult: 2}
    convolution_param {
        num_output: 20
        kernel_size: 5
        stride: 1
        weight_filler {
            type: "xavier"
        }
    }
}
....
```

- + Easy to deploy (C++)
- + A lot of pre-trained models (model zoo)
- Model as protobuf
- Hard to build new layers
- Hard to debug

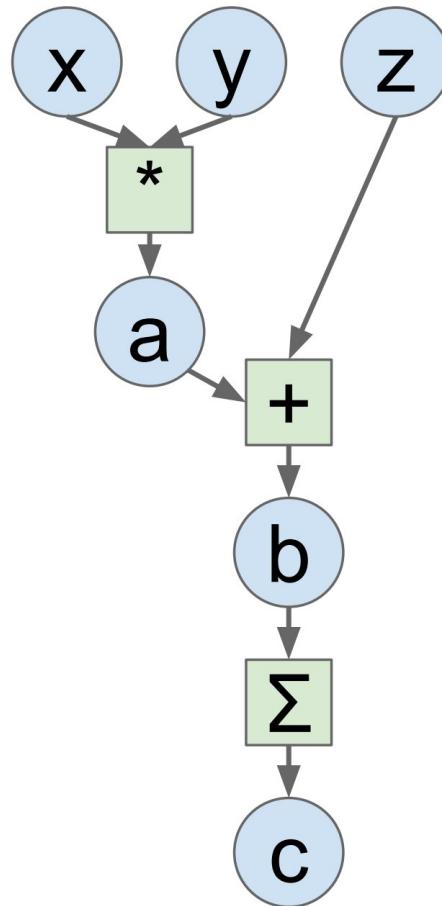
Still used in some legacy codebases

# Symbolic graphs

What will your CPU do  
when you write this?

```
a = x * y  
b = a + z  
c = np.sum(b)
```

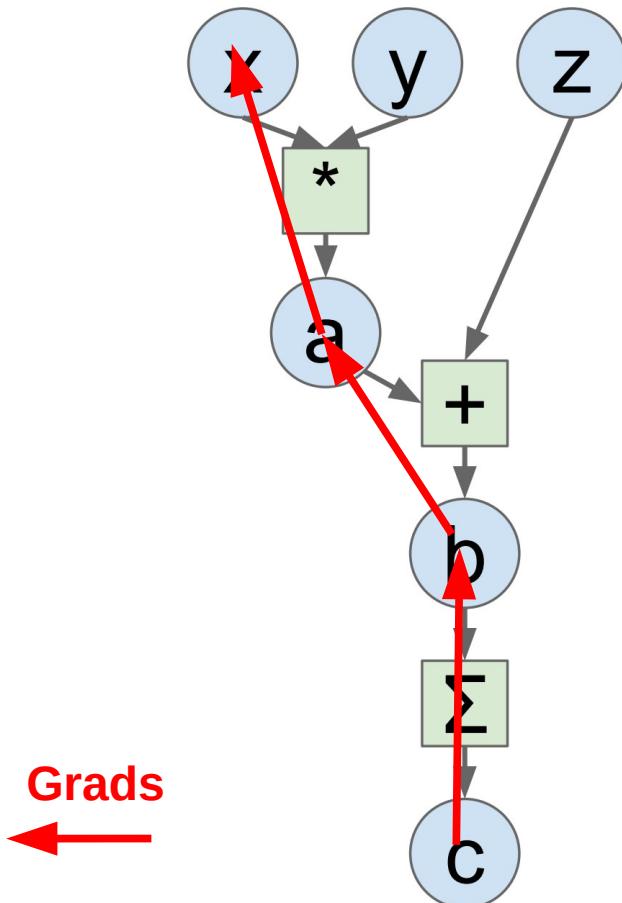
# Symbolic graphs



```
a = x * y  
b = a + z  
c = np.sum(b)
```

Idea: let's define  
this graph explicitly!

# Symbolic graphs



```
a = x * y  
b = a + z  
c = np.sum(b)
```

- + Automatic gradients!
- + Easy to build new layers
- + We can optimize the Graph
- Graph is static during training
- Need time to compile/optimize
- Hard to debug

# Symbolic graphs

## Static graph frameworks

- Purely static is legacy

Theano (deprecated)

TensorFlow (before 2.0)

- Static in modern frameworks  
`torch.jit.trace/script, compile`  
`jax.jit / tensorflow.function`



theano

# Dynamic graphs

Chainer, DyNet, Pytorch

A graph is created on the fly



```
from torch.autograd import Variable  
  
x = Variable(torch.randn(1, 10))  
prev_h = Variable(torch.randn(1, 20))  
W_h = Variable(torch.randn(20, 20))  
W_x = Variable(torch.randn(20, 10))
```

# Dynamic graphs

Chainer, DyNet, Pytorch



- + Can change graph on the fly
- + Can get value of any tensor at any time  
(easy debugging)
- Hard to optimize graphs  
(especially large graphs)
- Still early development

**Researchers love them!**

# Dynamic graphs



**Andrej Karpathy** @karpathy

Following

I've been using PyTorch a few months now and I've never felt better. I have more energy. My skin is clearer. My eye sight has improved.

**Researchers love them!**

# Advanced: GPU kernels w/o CUDA

<https://openai.com/index/triton/>

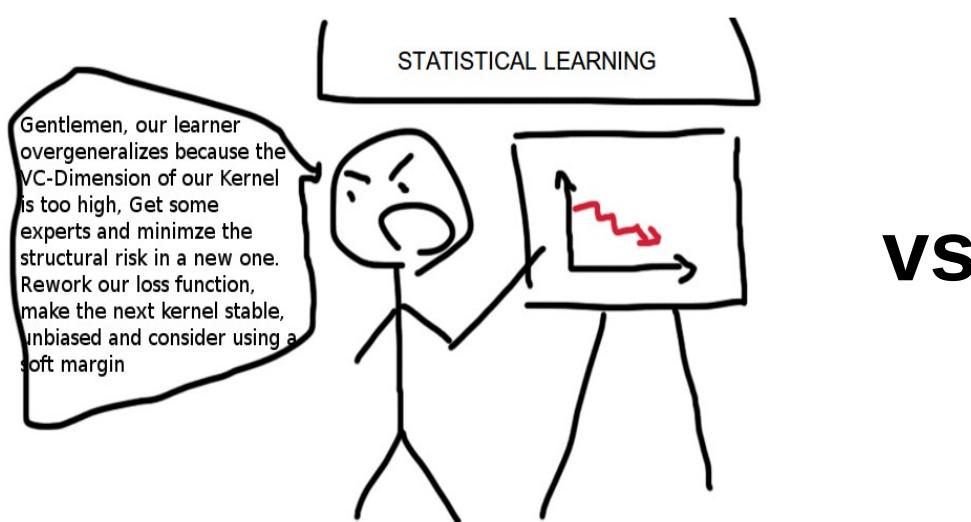
Python

```
1  @triton.jit
2  def matmul(A, B, C, M, N, K, stride_am, stride_ak,
3             stride_bk, stride_bn, stride_cm, stride_cn,
4             **META):
5     # extract metaparameters
6     BLOCK_M, GROUP_M = META['BLOCK_M'], META['GROUP_M']
7     BLOCK_N = META['BLOCK_N']
8     BLOCK_K = META['BLOCK_K']
9     # programs are grouped together to improve L2 hit rate
10    _pid_m = tl.program_id(0)
11    _pid_n = tl.program_id(1)
12    pid_m = _pid_m // GROUP_M
13    pid_n = (_pid_n * GROUP_M) + (_pid_m % GROUP_M)
14    # rm (resp. rn) denotes a range of indices
15    # for rows (resp. col) of C
16    rm = pid_m * BLOCK_M + tl.arange(0, BLOCK_M)
17    rn = pid_n * BLOCK_N + tl.arange(0, BLOCK_N)
18    # rk denotes a range of indices for columns
19    # (resp. rows) of A (resp. B)
20    rk = tl.arange(0, BLOCK_K)
21    # the memory addresses of elements in the first block of
22    # A and B can be computed using numpy-style broadcasting
23    A = A + (rm[:, None] * stride_am + rk[None, :] * stride_ak)
24    B = B + (rk[:, None] * stride_bk + rn[None, :] * stride_bn)
25    # initialize and iteratively update accumulator
26    acc = tl.zeros((BLOCK_M, BLOCK_N), dtype=tl.float32)
```

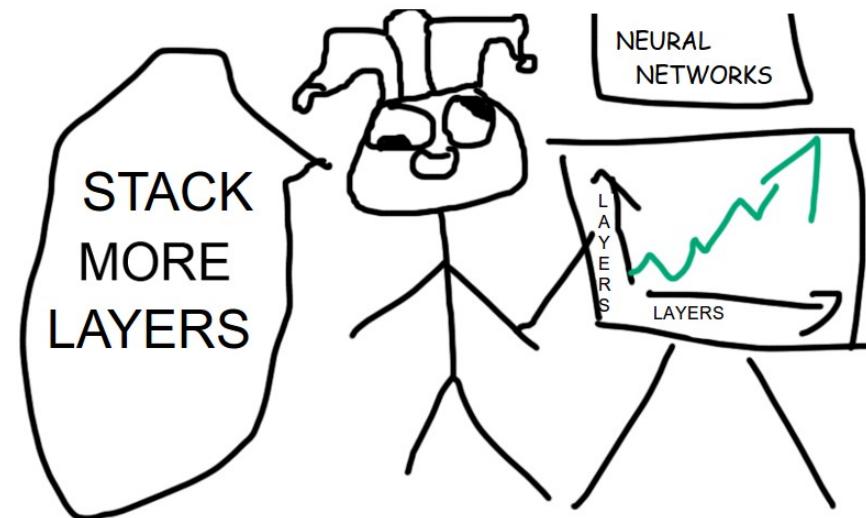
**[Short break, then practice, then the rest of the class]**

# Not magic!

Don't expect deep learning to solve all your problems for free. For it won't.



vs



<https://i.warosu.org/data/sci/img/0073/62/1435656449422.png>

# Not magic

## **Book of grudges**

- No core theory
  - Relies on intuitive reasoning

# Not magic

## **Book of grudges**

- No core theory
  - Relies on intuitive reasoning
- Needs tons of data
  - You need either large dataset or heavy wizardry

# Not magic

## Book of grudges

- No core theory
  - Relies on intuitive reasoning
- Needs tons of data
  - You need either large dataset or heavy wizardry
- Computationally heavy
  - Running on mobiles/embedded is a challenge

# Not magic

## Book of grudges

- No core theory
  - Relies on intuitive reasoning
- Needs tons of data
  - You need either large dataset or heavy wizardry
- Computationally heavy
  - Running on mobiles/embedded is a challenge
- Pathologically overhyped
  - People expect of it to make wonders

Deep learning is a language

Deep learning is a language  
in which you can hint your model  
on what you want it to learn

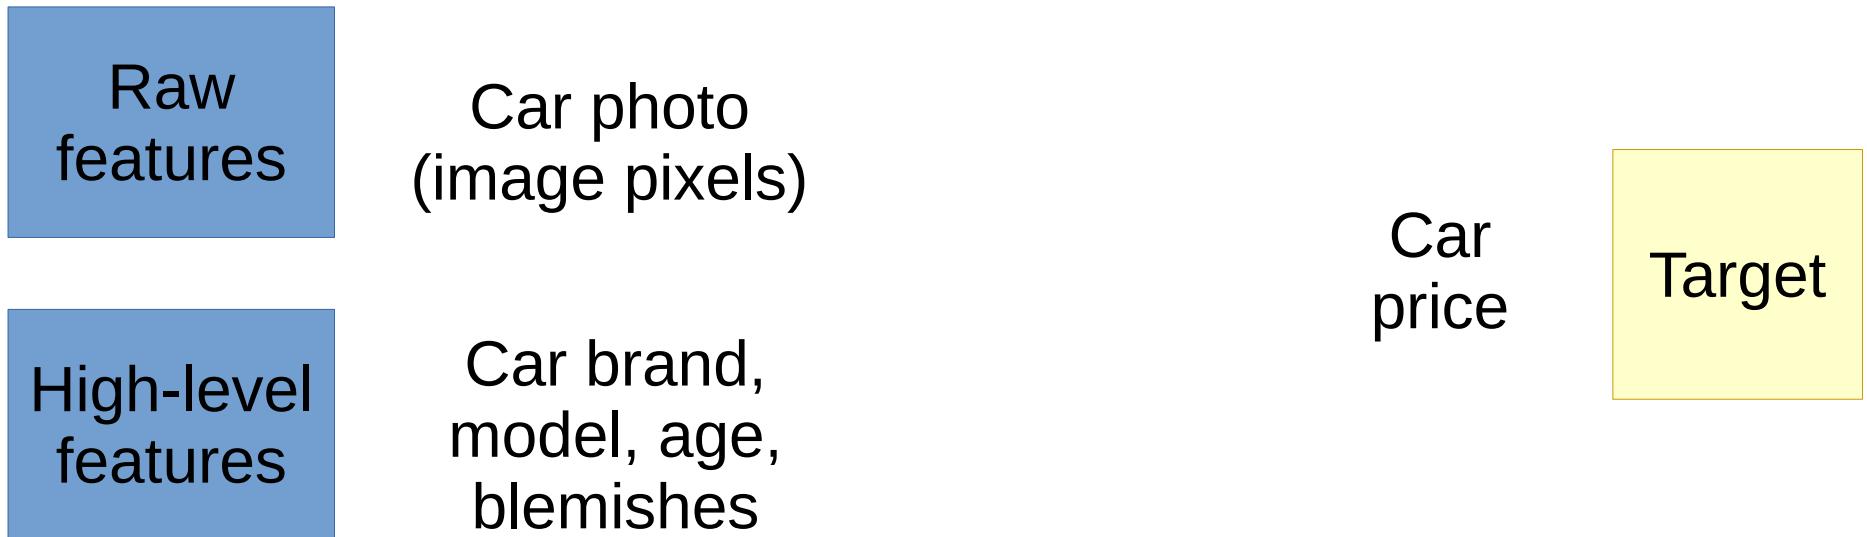
# Deep learning is a language

Say, you train classifier on two sets of features



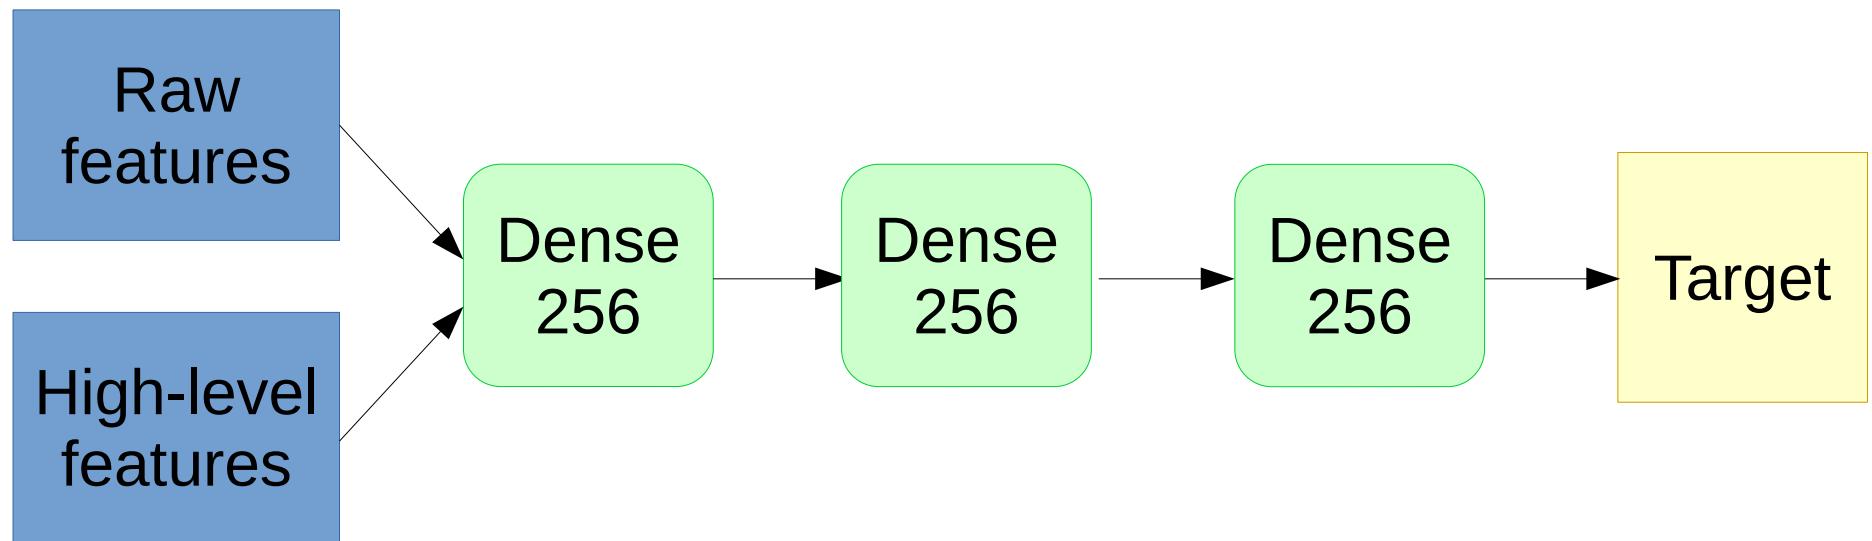
# Deep learning is a language

Say, you train classifier on two sets of features



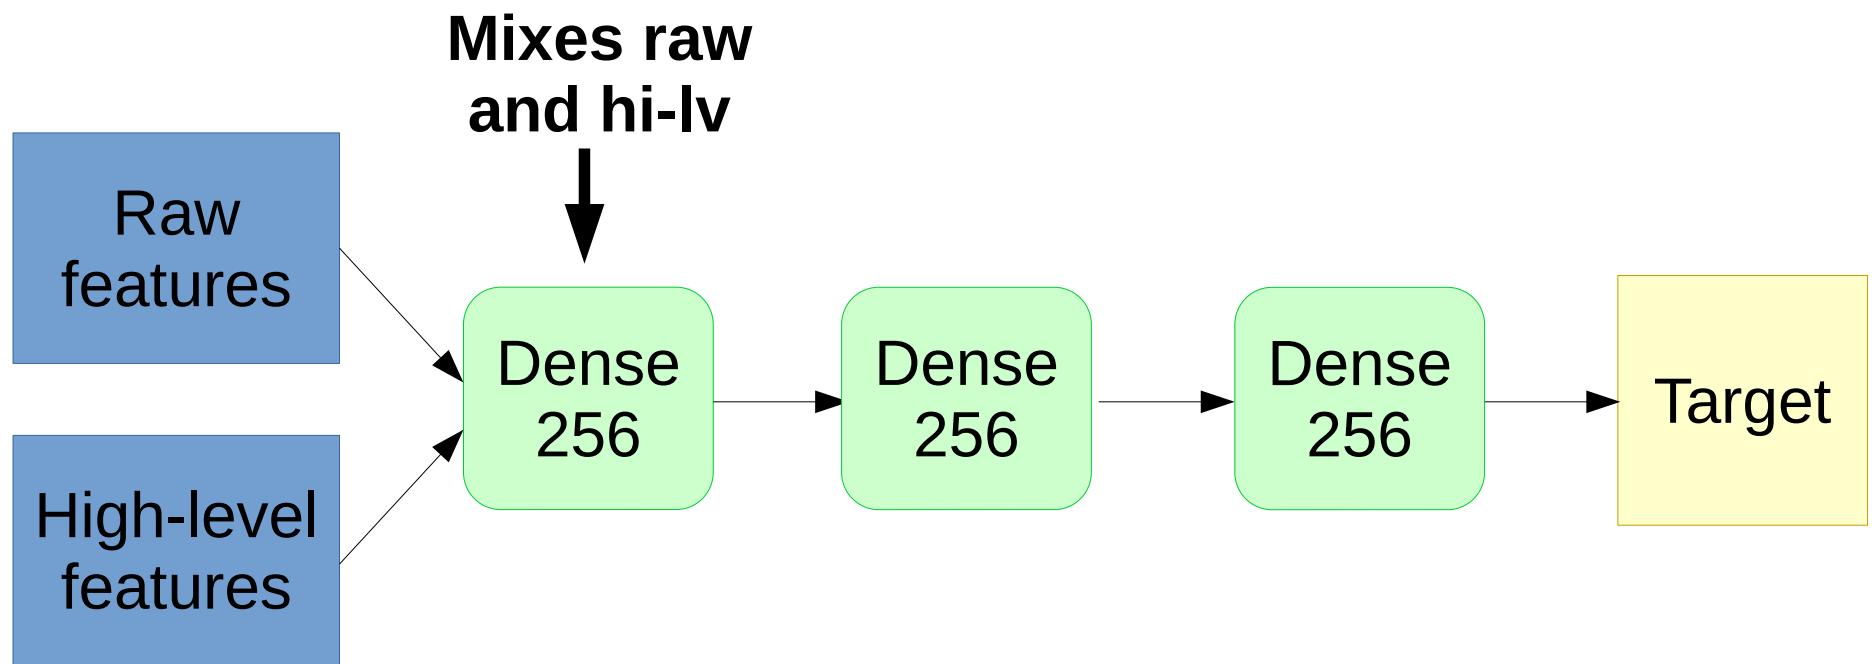
# Deep learning is a language

Naive approach



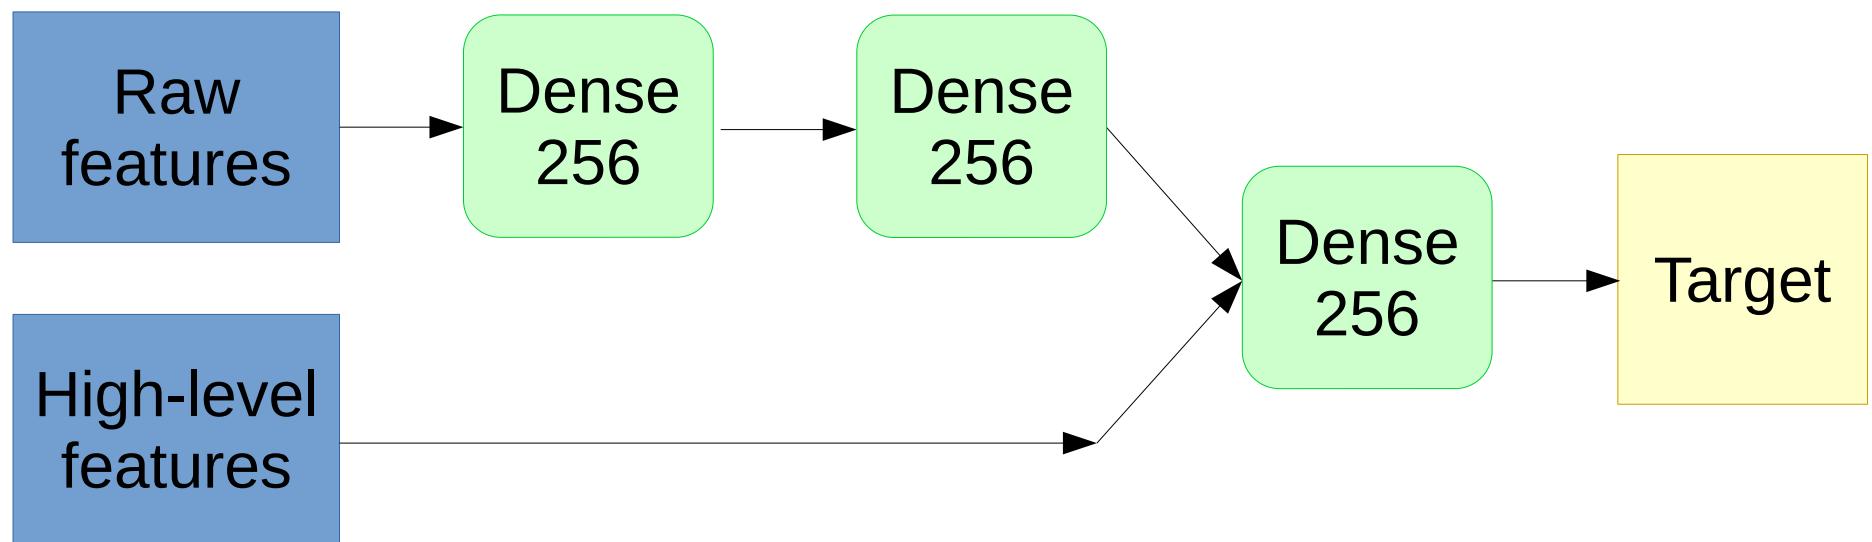
# Deep learning is a language

Naive approach



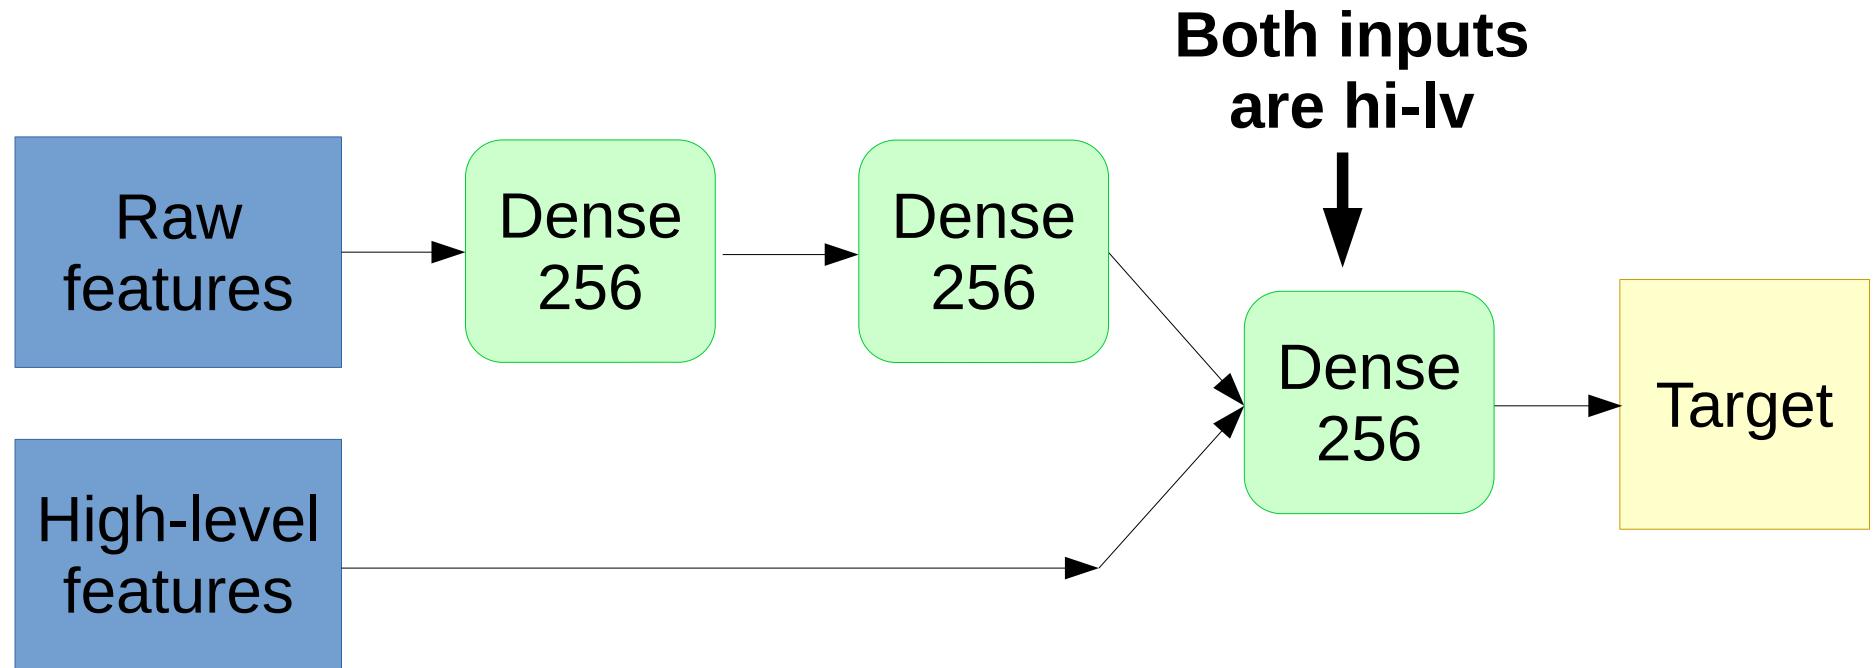
# Deep learning is a language

Less naïve approach



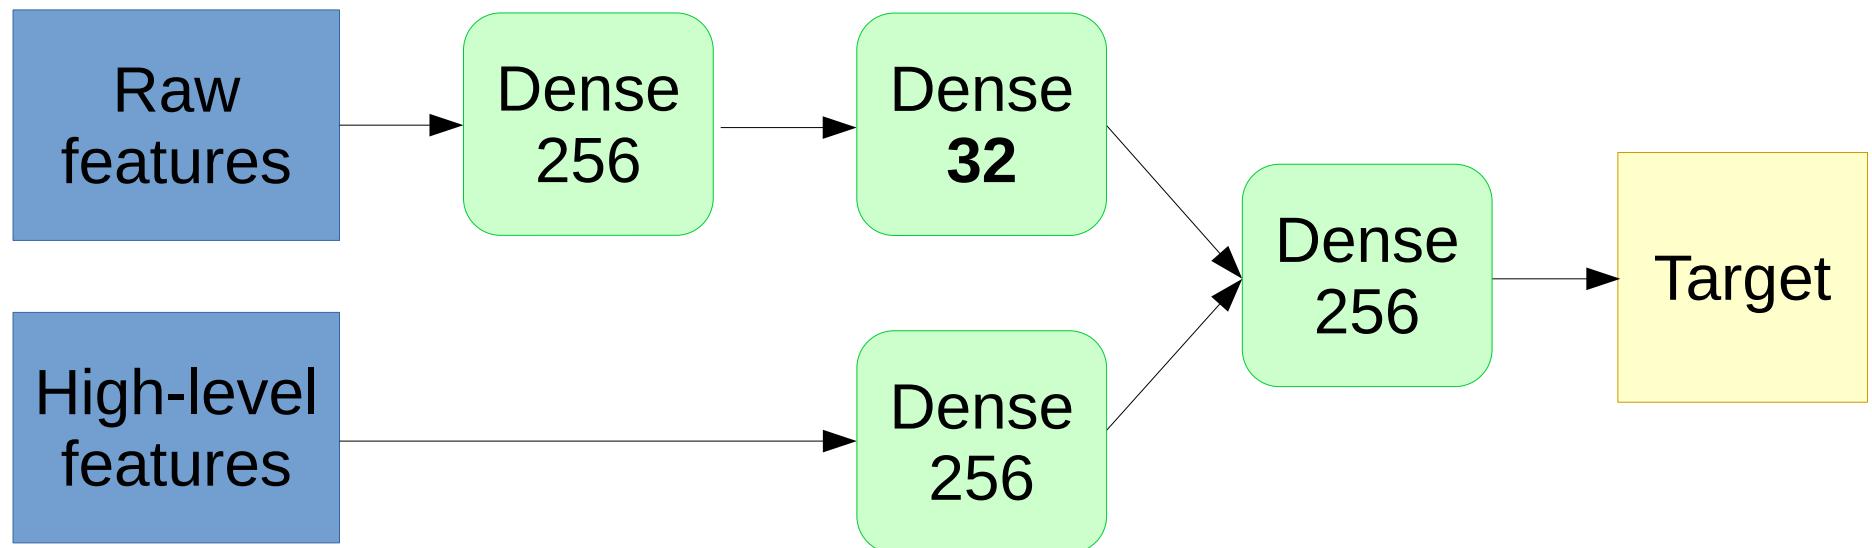
# Deep learning is a language

Less naïve approach



# Deep learning is a language

“Image features should be less important”  
*if that's what you want to say*



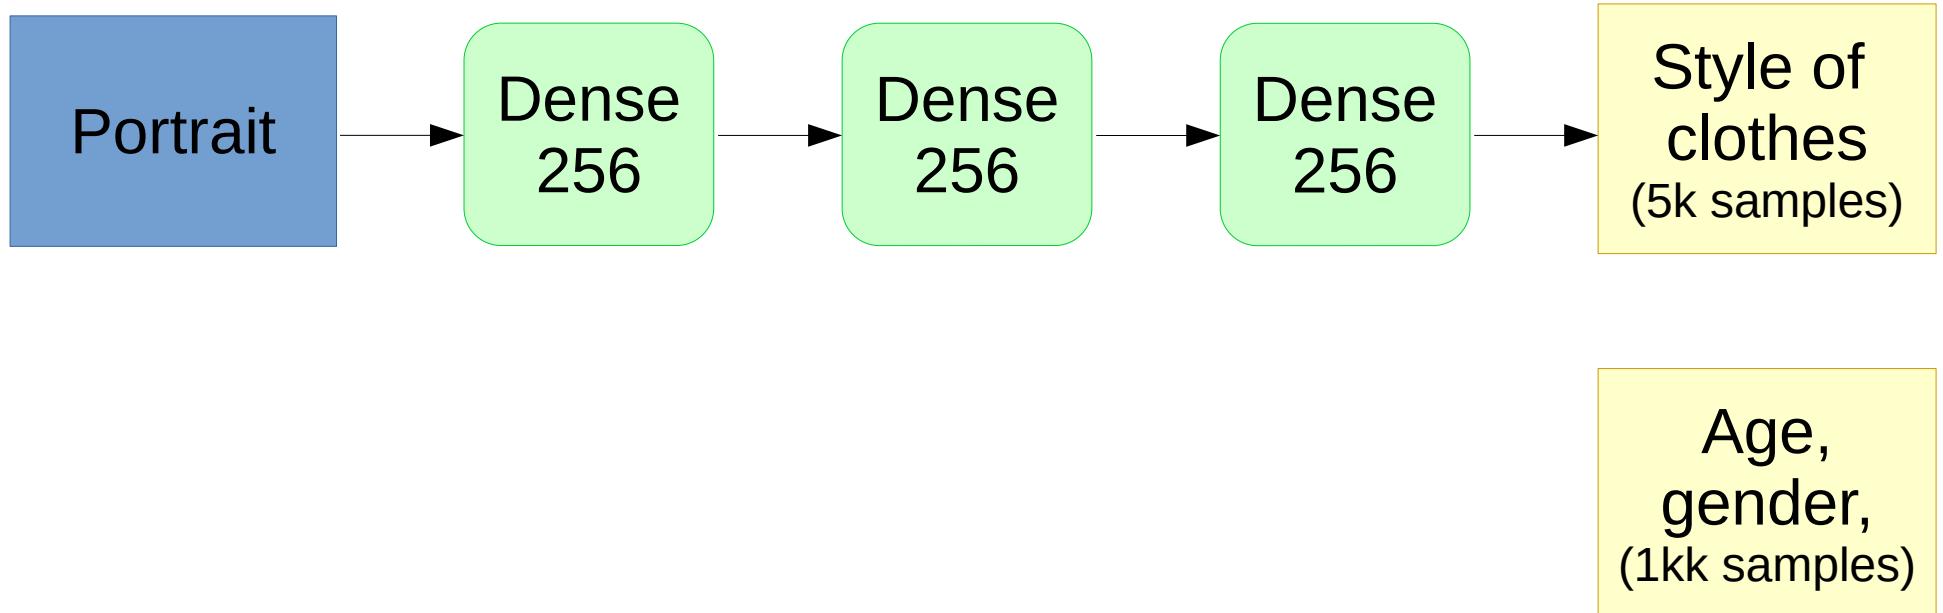
# Deep learning is a language

You have a small dataset



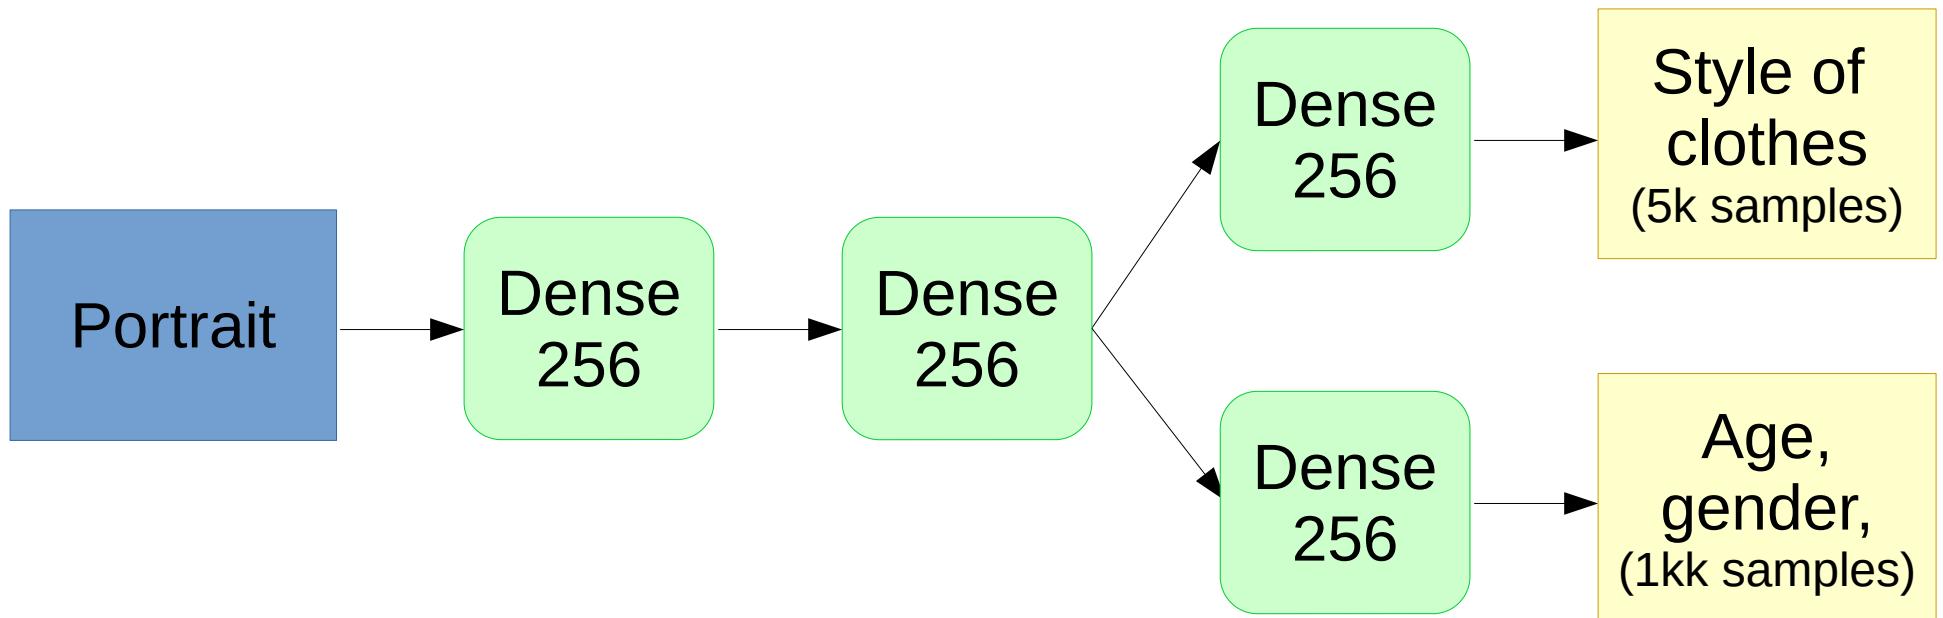
# Deep learning is a language

You have a small dataset  
and a larger dataset with similar task



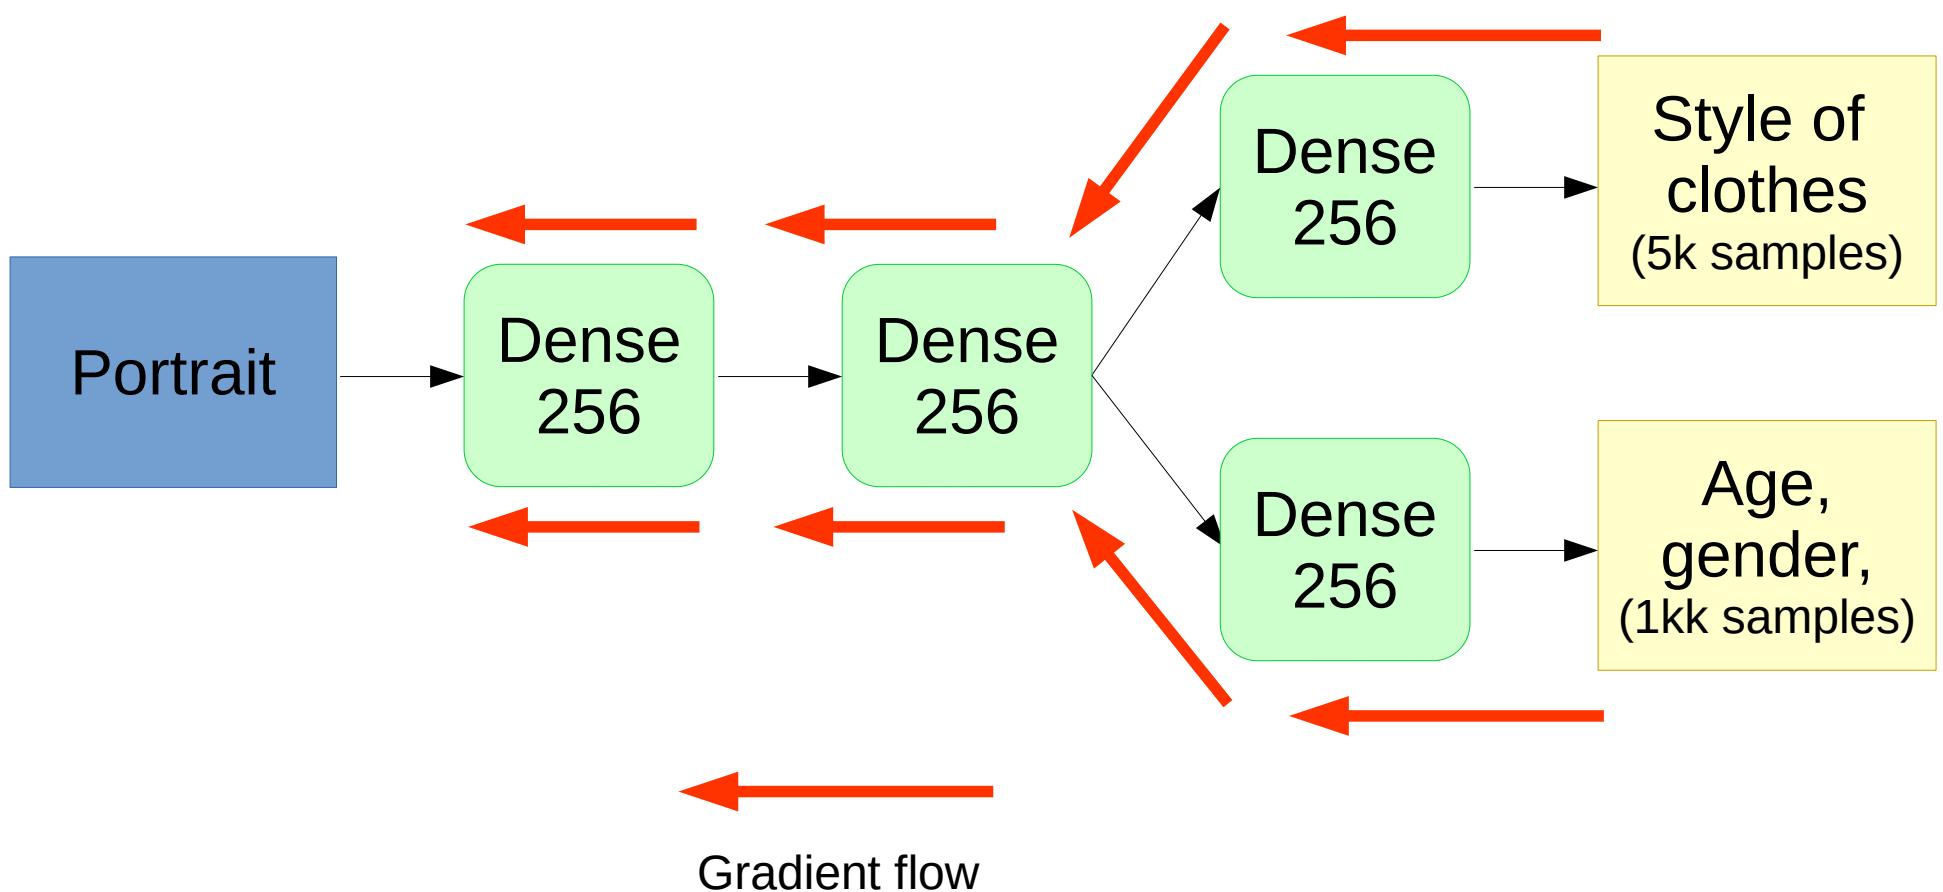
# Deep learning is a language

You have a small dataset  
and a larger dataset with similar task



# Deep learning is a language

I want to learn features for style classification  
that also help determine age & gender



# Deep learning is a language

For images:

- “I want to classify cats regardless where they are”
- “A cat shifted by 3 pixels is still a cat”

For texts:

- “People read and write texts left to right”

In general:

- “I don't want model to trust single feature too much”
- “I want my features to be sparse”

Let's see a few more “words”

# Regularization

- Neural networks overfit like nothing else.  
Gotta regularize!
- We can use L1/L2 like usual, but there's more!

# Regularization

- Dropout:

“I don't trust my network to trust any single neuron too much”
- Idea:

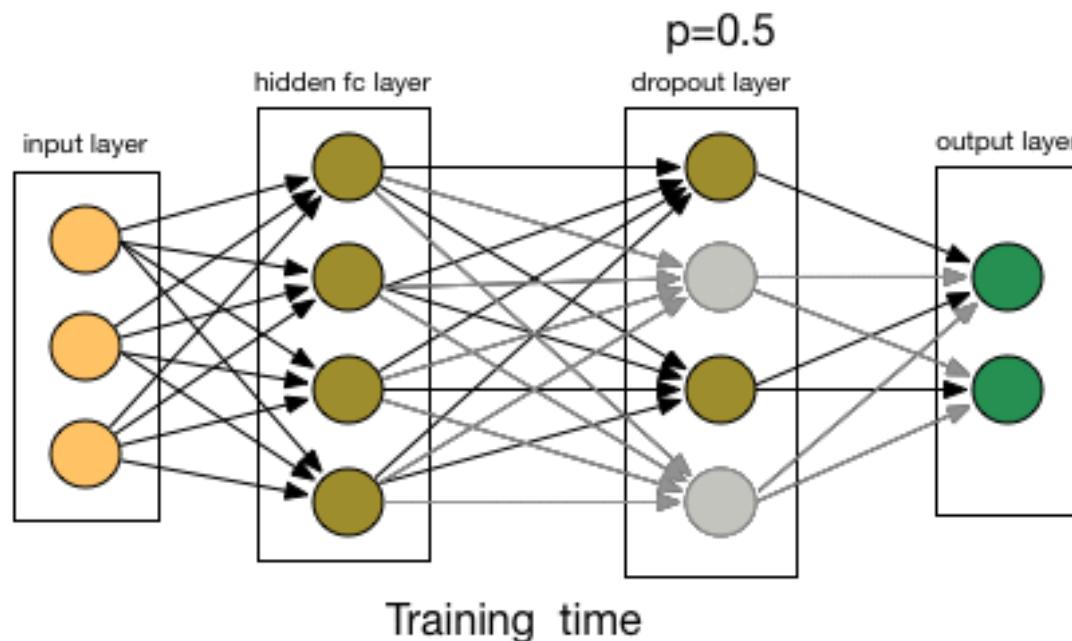
At training time, with probability  $p$   
multiply neurons by zero!

  - Scale up the remaining neurons to keep average the same

# Regularization

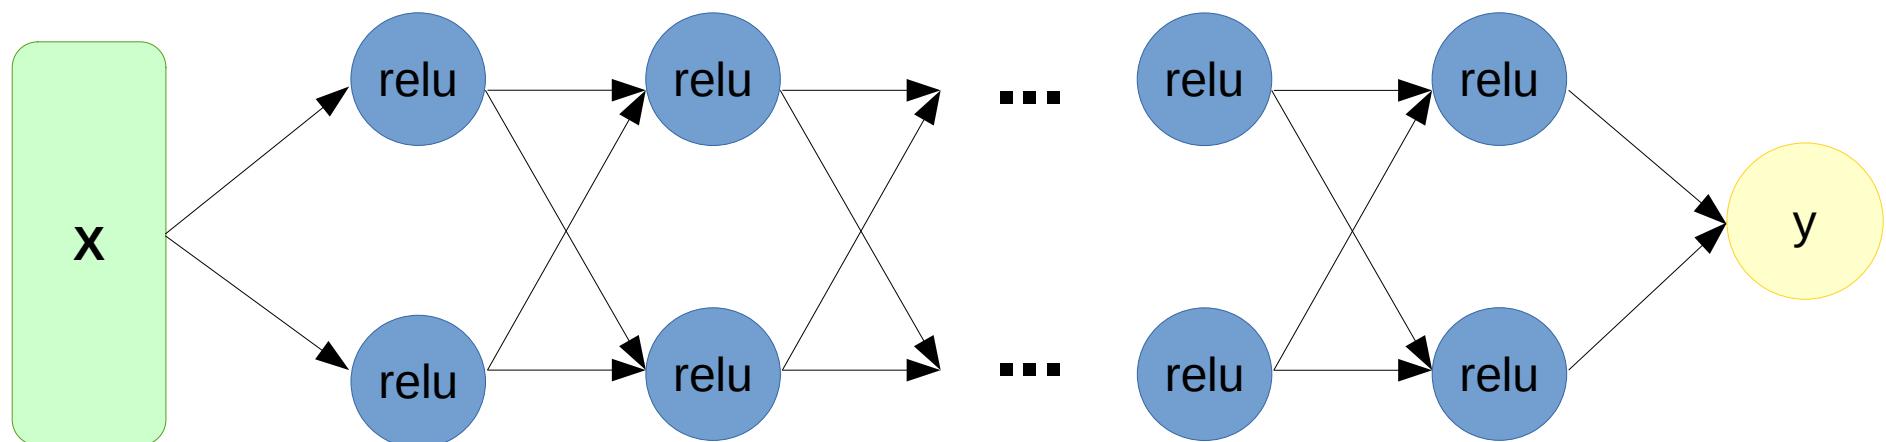
- Dropout:

“I don't trust my network to trust any single neuron too much”



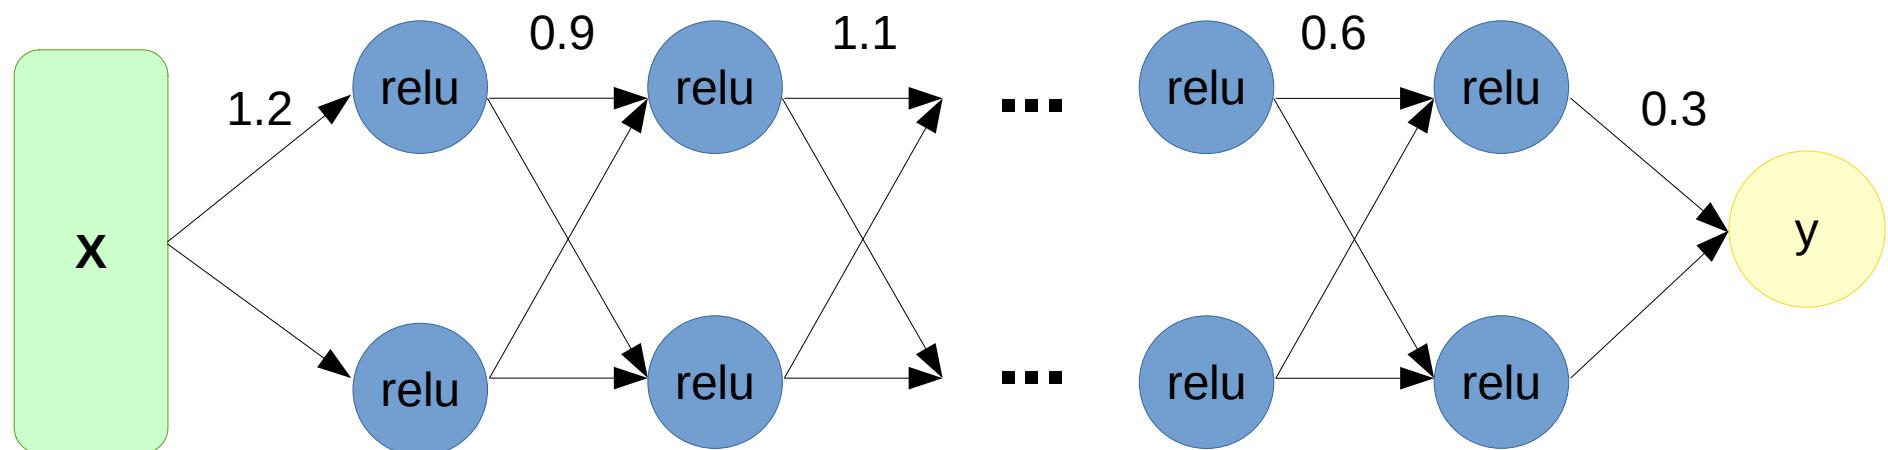
# The problem with deep networks

- Imagine a 100-layer network with ReLU



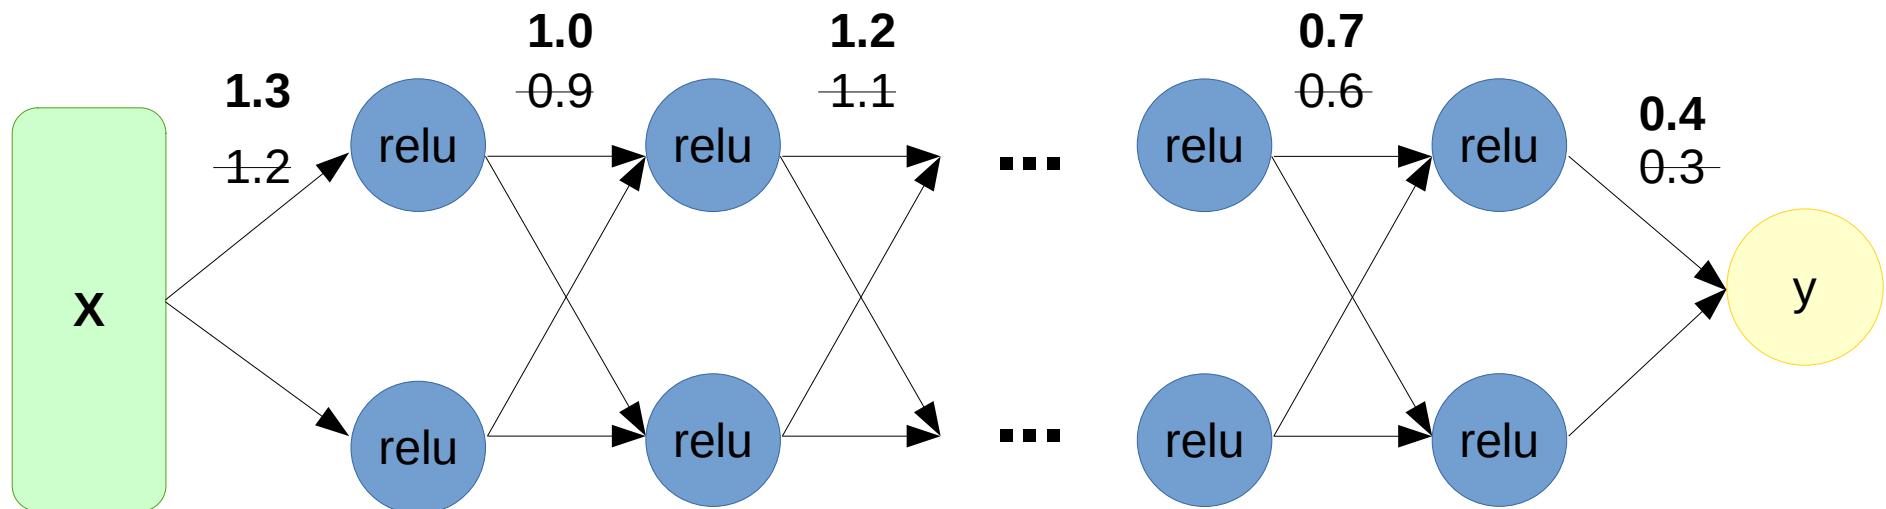
# The problem with deep networks

- Imagine a 100-layer network with ReLU



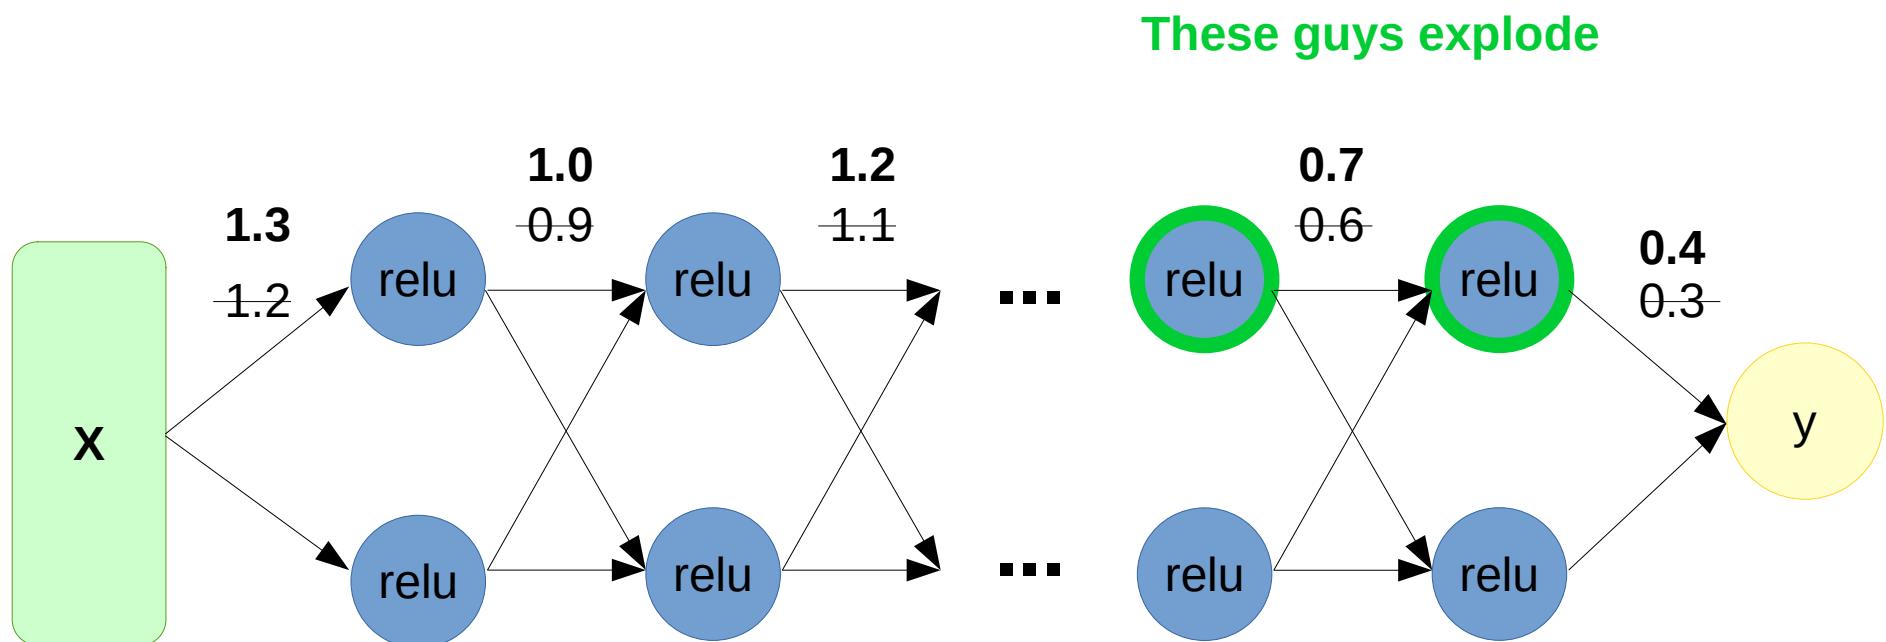
# The problem with deep networks

- Imagine a 100-layer network with ReLU
- Single gradient step...



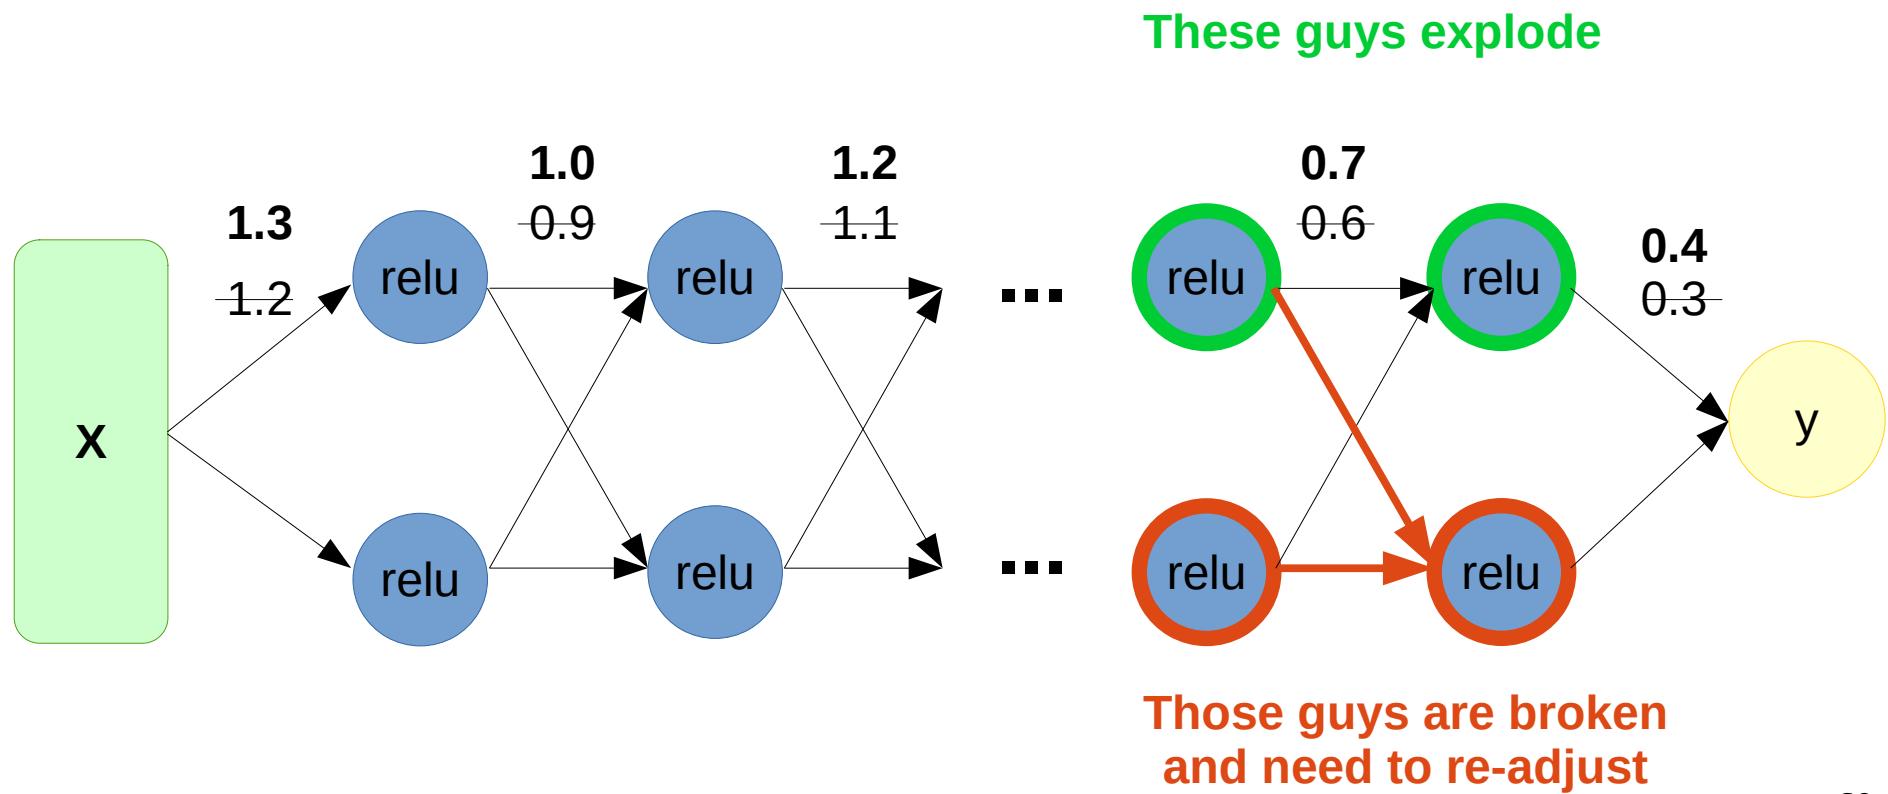
# The problem with deep networks

- Imagine a 100-layer network with ReLU
- Single gradient step...



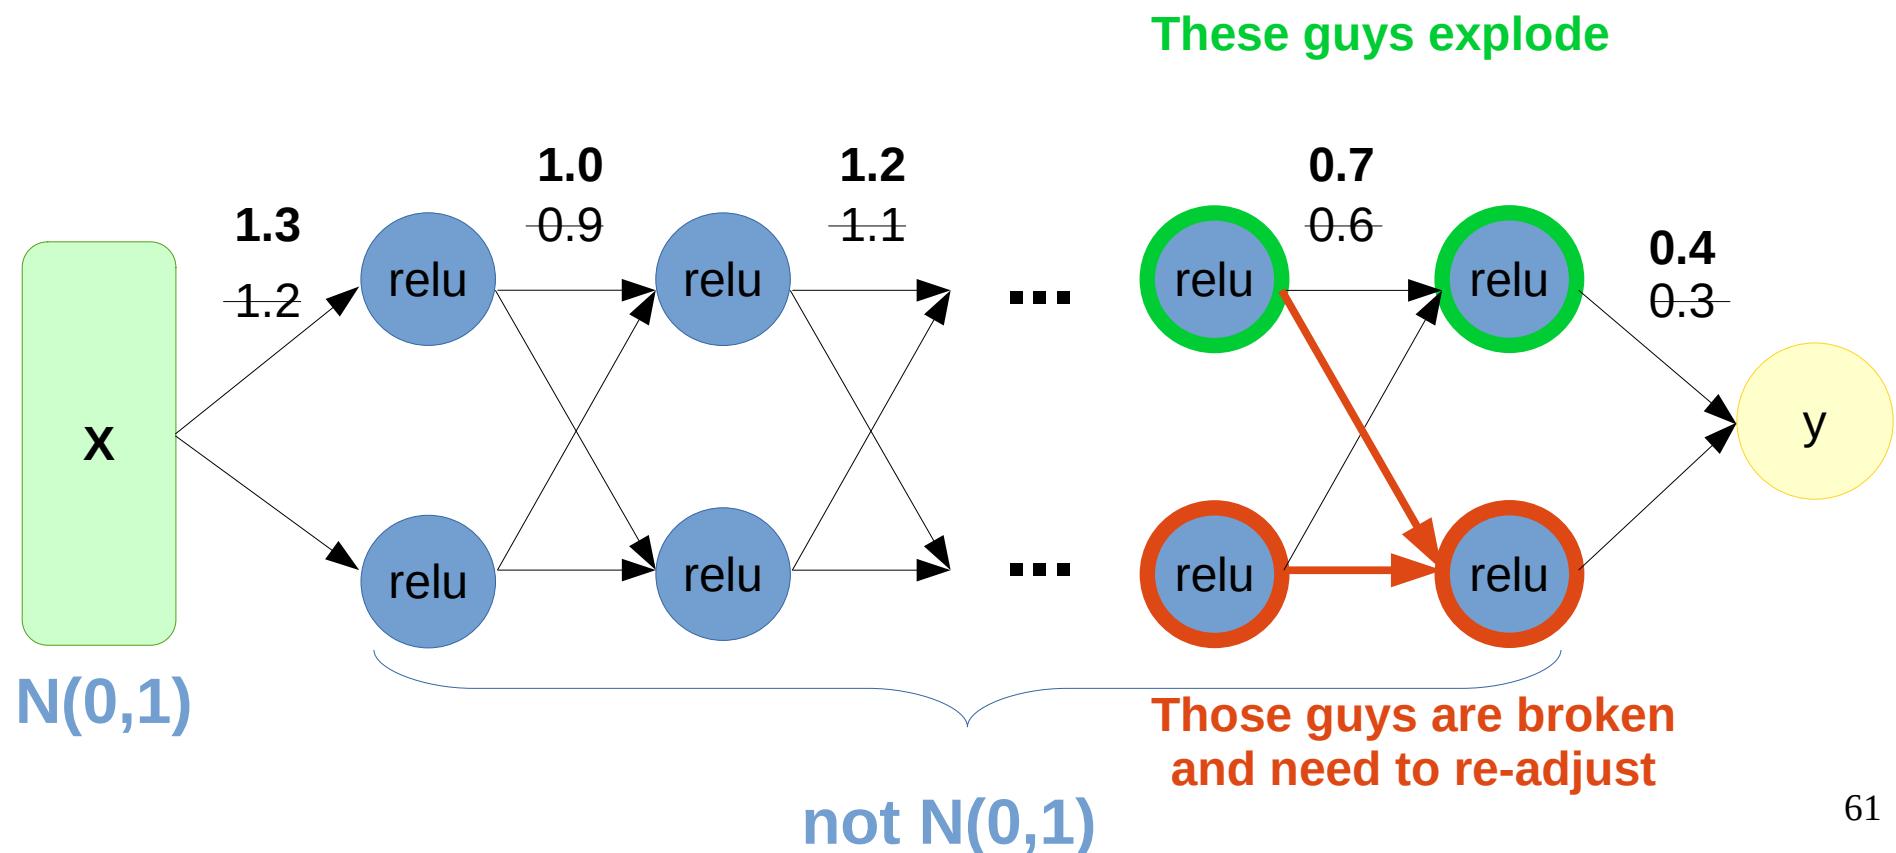
# The problem with deep networks

- Imagine a 100-layer network with ReLU
- Single gradient step...



# The problem with deep networks

- Imagine a 100-layer network with ReLU
- Single gradient step...



# Batch normalization

TL;DR:

- It's usually a good idea to normalize linear model inputs
  - (c) Every machine learning lecturer, ever

# Batch normalization

Idea:

- We normalize activation of a hidden layer  
(zero mean unit variance)

$$h_i = \frac{h_i - \mu_i}{\sqrt{\sigma_i^2}}$$

- Update  $\mu_i, \sigma_i^2$  with moving average while training

$$\mu_i := \alpha \cdot \text{mean}_{batch} + (1 - \alpha) \cdot \mu_i$$

$$\sigma_i^2 := \alpha \cdot \text{variance}_{batch} + (1 - \alpha) \cdot \sigma_i^2$$

# Batch normalization

Idea:

- We normalize activation of a hidden layer  
(zero mean unit variance)

$$h_i = \frac{h_i - \mu_i}{\sqrt{\sigma_i^2}}$$

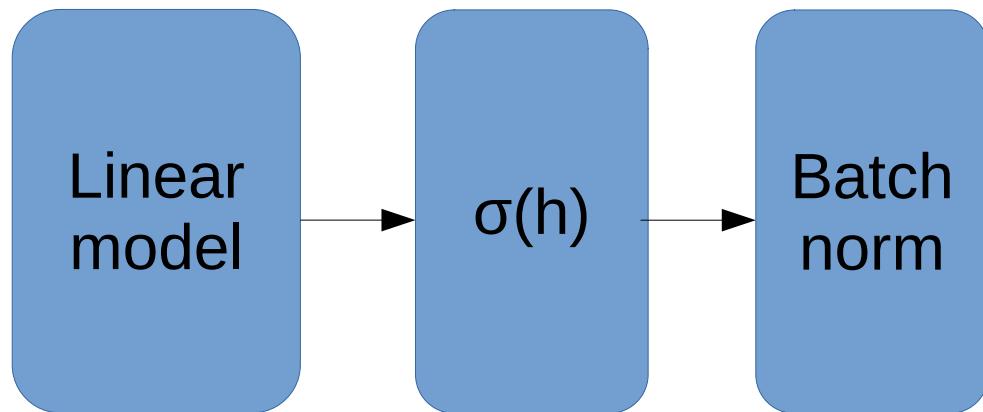
i stands for i-th neuron

- Update  $\mu_i, \sigma_i^2$  with moving average while training

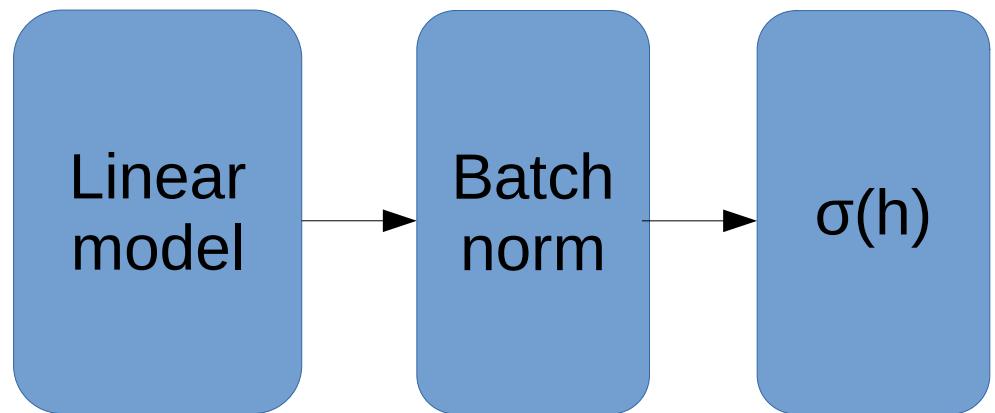
$$\mu_i := \alpha \cdot \text{mean}_{batch} + (1 - \alpha) \cdot \mu_i$$

$$\sigma_i^2 := \alpha \cdot \text{variance}_{batch} + (1 - \alpha) \cdot \sigma_i^2$$

# Batch normalization



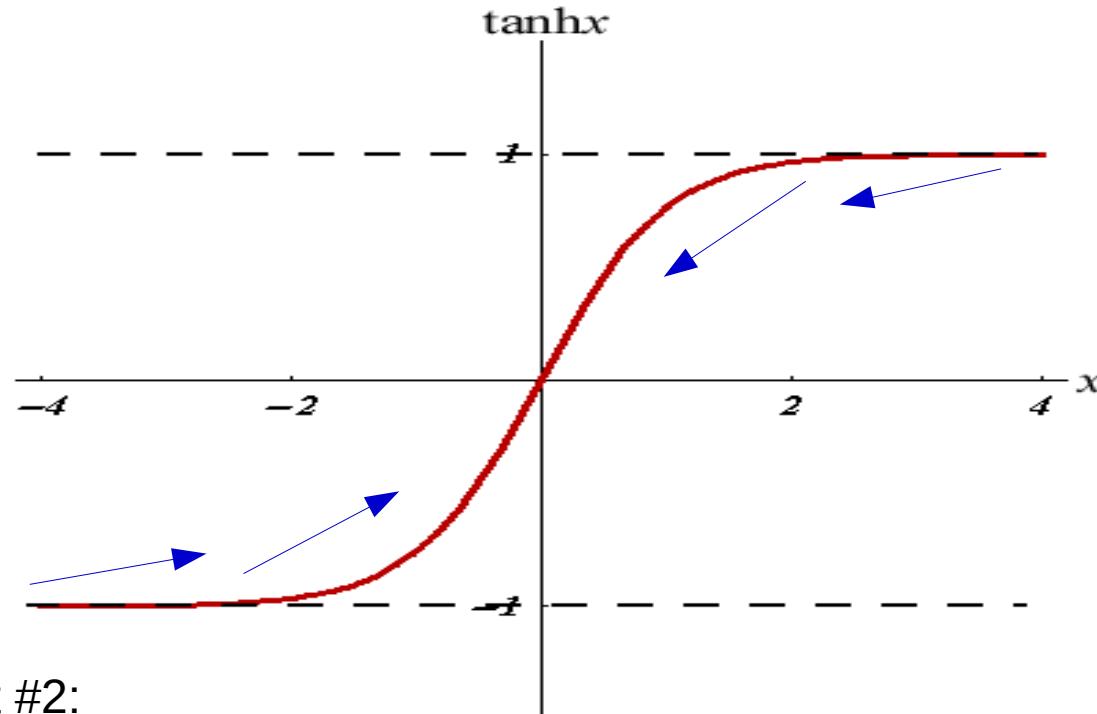
VS



# Batch normalization

Good side effect #1:

- Vanishing gradient less a problem for sigmoid-like nonlinearities



Good side effect #2:

- We no longer need to train bias ( $+b$  term in  $Wx+b$ )

# Weight normalization

Same problem, different solution

- Learn separate “direction”  $w$  and “length”  $l$

$$\hat{w} \stackrel{\text{def}}{=} \frac{w}{\|w\|} \cdot l$$

- Much simpler, but requires good init

# More normalization

## Layer/Instance normalization

- Like batchnorm, but normalizes over different axes

## Normprop

- A special training algorithm

## Self-normalizing neural networks (SELU)

# More normalization

## Layer/Instance normalization

- Like batchnorm, but normalizes over different axes

## Normprop

- A special training algorithm

## Self-normalizing neural networks (SELU)

# Architectures: residual network

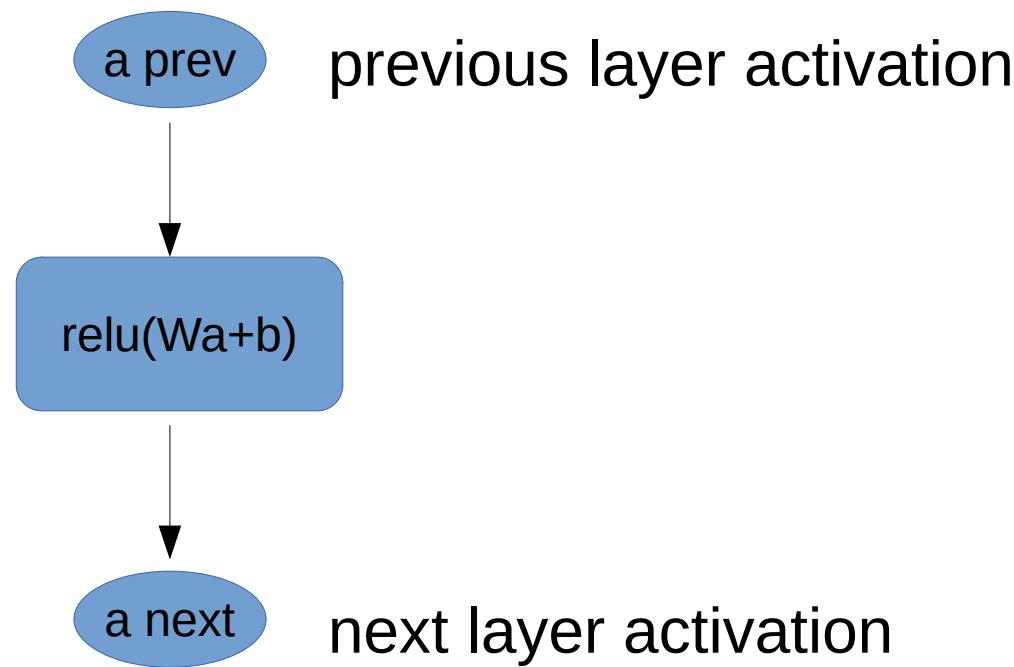
Problem: very deep networks are hard to train

Gradients w.r.t. first layers are volatile

# Architectures: residual network

Idea: let's create a shortcut for gradients

Normal layer

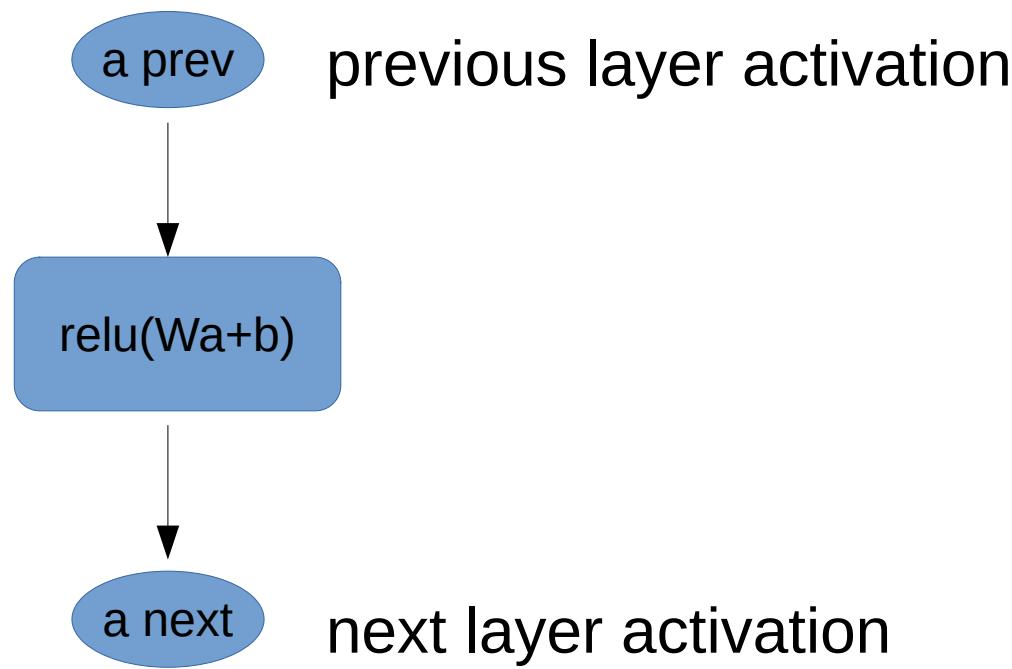


$$f_{w,b}(x) = \text{relu}(W \cdot a + b)$$

# Architectures: residual network

Idea: let's create a shortcut for gradients

Normal layer



$$f_{w,b}(x) = \text{relu}(W \cdot a + b)$$

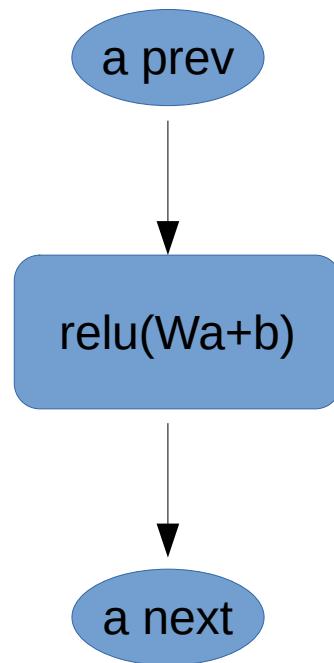
$$\nabla f_{w,b}(x) = \nabla \text{relu}(W \cdot a + b)$$

**Gradients can vanish if  $\text{relu} < 0$**  <sup>72</sup>

# Architectures: residual network

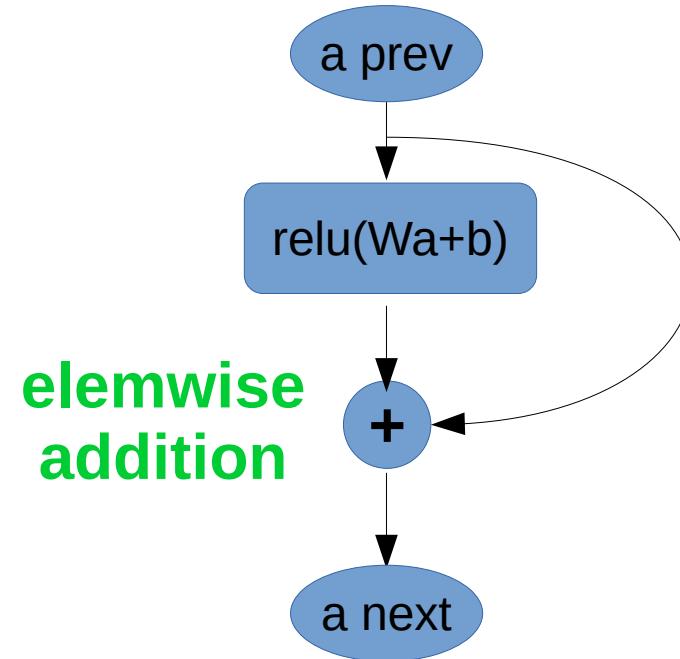
Idea: let's create a shortcut for gradients

Normal layer



$$f_{w,b}(x) = \text{relu}(W \cdot a + b)$$

Residual layer



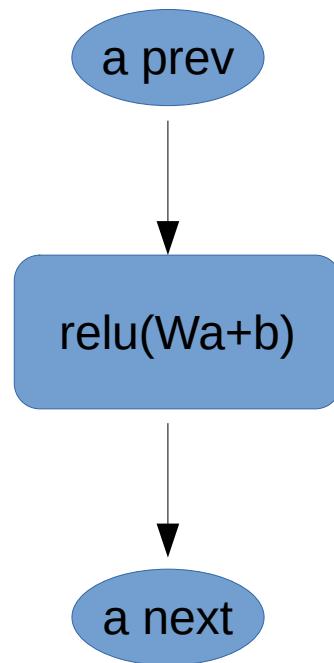
elemwise  
addition

$$f_{w,b}(x) = \text{relu}(W \cdot a + b) + X$$

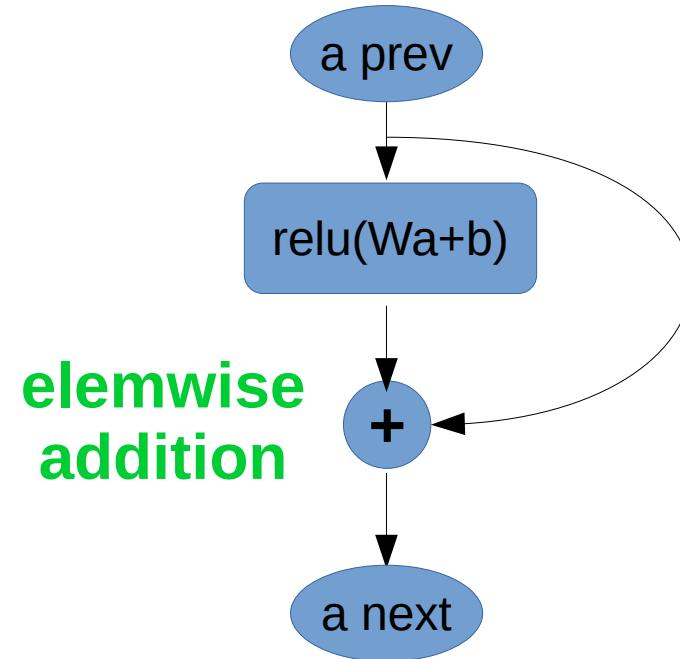
# Architectures: residual network

Idea: let's create a shortcut for gradients

Normal layer



Residual layer



$$f_{w,b}(x) = \text{relu}(W \cdot a + b)$$

$$\nabla f_{w,b}(x) = \nabla \text{relu}(W \cdot a + b)$$

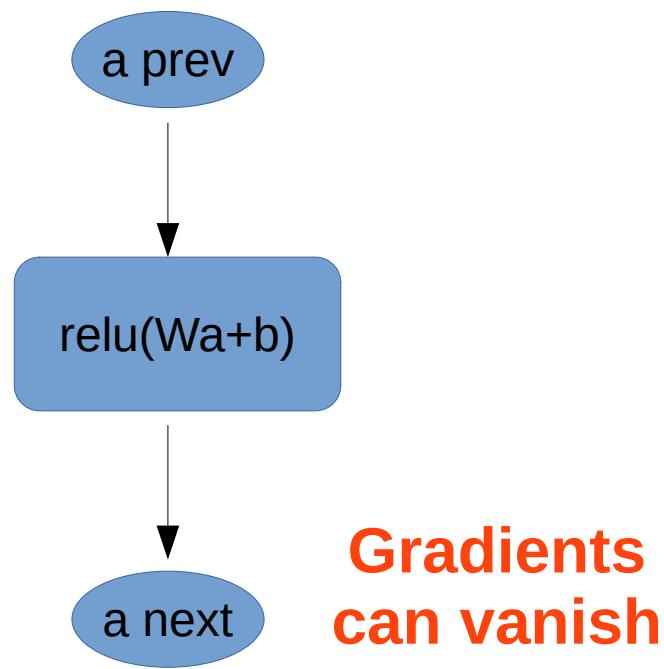
$$f_{w,b}(x) = \text{relu}(W \cdot a + b) + X$$

???

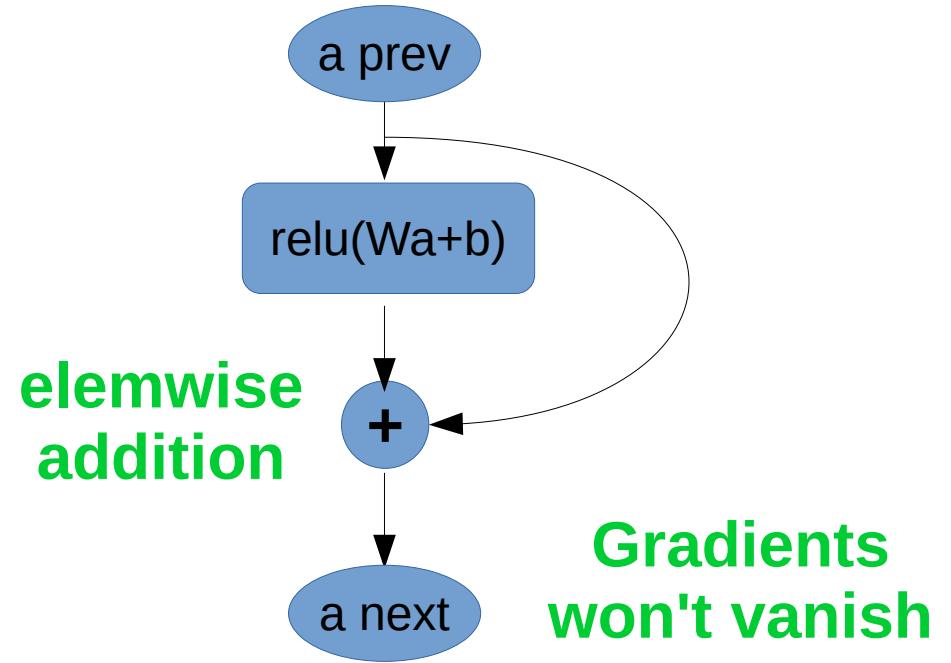
# Architectures: residual network

Idea: let's create a shortcut for gradients

Normal layer



Residual layer



$$f_{w,b}(x) = \text{relu}(W \cdot a + b)$$

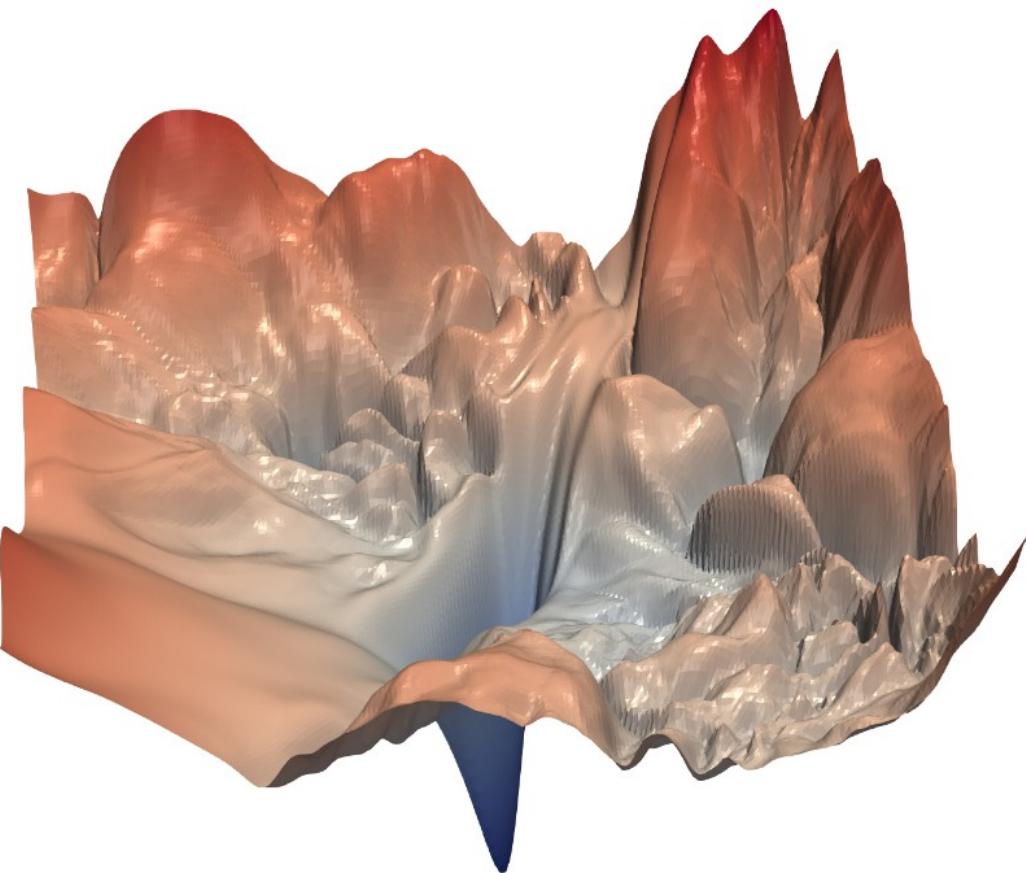
$$\nabla f_{w,b}(x) = \nabla \text{relu}(W \cdot a + b)$$

$$f_{w,b}(x) = \text{relu}(W \cdot a + b) + X$$

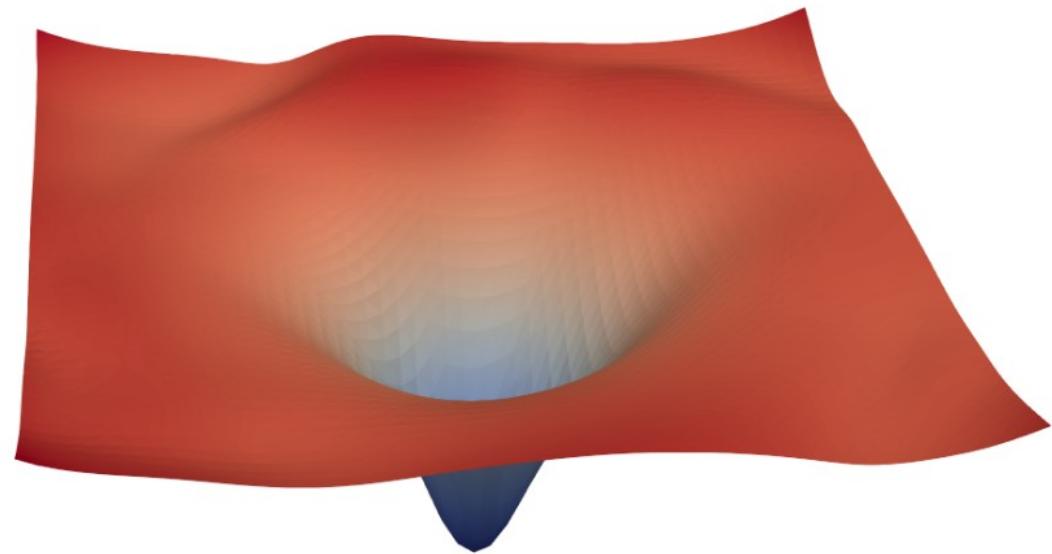
$$\nabla f_{w,b}(x) = \nabla \text{relu}(W \cdot a + b) + \vec{1}$$

# Loss Surfaces

Idea: <https://arxiv.org/abs/1712.09913>



(a) without skip connections



(b) with skip connections

# TL;DR today

Didn't like writing backward code?

# TL;DR today

Didn't like writing backward code?

Autograd (pytorch, jax, tf, etc)

Multiple types of input data?

e.g. image + tabular

# TL;DR today

Didn't like writing backward code?

Autograd (pytorch, jax, tf, etc)

Multiple types of input data?  
e.g. image + tabular

Concat / add branches

Model overfits too quickly?

# TL;DR today

Didn't like writing backward code?

Autograd (pytorch, jax, tf, etc)

Multiple types of input data?  
e.g. image + tabular

Concat / add branches

Model overfits too quickly?

Dropout or similar

Many layers hard to train?

# TL;DR today

Didn't like writing backward code?

Autograd (pytorch, jax, tf, etc)

Multiple types of input data?  
e.g. image + tabular

Concat / add branches

Model overfits too quickly?

Dropout or similar

Many layers hard to train?

BatchNorm, Residual connections

# Nuff

## Coding time!

